

Forecasting Hotel Booking Cancellations.

The hospitality industry thrives in tourist hotspots, but one of the significant challenges it faces is booking cancellations. Last-minute or unexpected cancellations can result in substantial revenue losses for hotels. When a booking is canceled, the hotel not only loses the expected income but also misses the opportunity to rebook the room for the same period.

Accurately predicting the likelihood of a booking cancellation in advance can help hotel owners mitigate these losses.

The aim of this project is to forecast room cancellations using historical booking data with the highest possible accuracy

Importing necessary libraries

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import sklearn

import warnings
warnings.filterwarnings('ignore')
```

Getting the data

```
df =
pd.read_csv("/kaggle/input/hotel-booking-demand/hotel_bookings.csv")
df.head()
```

	hotel	is_canceled	lead_time	arrival_date_year
0	Resort Hotel	0	342	2015
1	Resort Hotel	0	737	2015
2	Resort Hotel	0	7	2015
3	Resort Hotel	0	13	2015
4	Resort Hotel	0	14	2015

	arrival_date_week_number	arrival_date_day_of_month
0	27	1
1	27	1
2	27	1

3	27	1
4	27	1

	stays_in_weekend_nights	stays_in_week_nights	adults	...	
deposit_type \					
0	0	0	2	...	No
Deposit					
1	0	0	2	...	No
Deposit					
2	0	1	1	...	No
Deposit					
3	0	1	1	...	No
Deposit					
4	0	2	2	...	No
Deposit					

	agent	company	days_in_waiting_list	customer_type	adr	\
0	NaN	NaN	0	Transient	0.0	
1	NaN	NaN	0	Transient	0.0	
2	NaN	NaN	0	Transient	75.0	
3	304.0	NaN	0	Transient	75.0	
4	240.0	NaN	0	Transient	98.0	

	required_car_parking_spaces	total_of_special_requests
reservation_status \		
0	0	0
Check-Out		
1	0	0
Check-Out		
2	0	0
Check-Out		
3	0	0
Check-Out		
4	0	1
Check-Out		

	reservation_status_date
0	2015-07-01
1	2015-07-01
2	2015-07-02
3	2015-07-02
4	2015-07-03

[5 rows x 32 columns]

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 119390 entries, 0 to 119389
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	hotel	119390 non-null	object
1	is_canceled	119390 non-null	int64
2	lead_time	119390 non-null	int64
3	arrival_date_year	119390 non-null	int64
4	arrival_date_month	119390 non-null	object
5	arrival_date_week_number	119390 non-null	int64
6	arrival_date_day_of_month	119390 non-null	int64
7	stays_in_weekend_nights	119390 non-null	int64
8	stays_in_week_nights	119390 non-null	int64
9	adults	119390 non-null	int64
10	children	119386 non-null	float64
11	babies	119390 non-null	int64
12	meal	119390 non-null	object
13	country	118902 non-null	object
14	market_segment	119390 non-null	object
15	distribution_channel	119390 non-null	object
16	is_repeated_guest	119390 non-null	int64
17	previous_cancellations	119390 non-null	int64
18	previous_bookings_not_canceled	119390 non-null	int64
19	reserved_room_type	119390 non-null	object
20	assigned_room_type	119390 non-null	object
21	booking_changes	119390 non-null	int64
22	deposit_type	119390 non-null	object
23	agent	103050 non-null	float64
24	company	6797 non-null	float64
25	days_in_waiting_list	119390 non-null	int64
26	customer_type	119390 non-null	object
27	adr	119390 non-null	float64
28	required_car_parking_spaces	119390 non-null	int64
29	total_of_special_requests	119390 non-null	int64
30	reservation_status	119390 non-null	object
31	reservation_status_date	119390 non-null	object

dtypes: float64(4), int64(16), object(12)
memory usage: 29.1+ MB

- There are 119390 records in total and 32 columns.
- There are missing values in a few columns

Exploratory Data Analysis and Data Pre-processing

```
df.isnull().sum()/len(df) * 100
```

hotel	0.000000
is_canceled	0.000000
lead_time	0.000000
arrival_date_year	0.000000
arrival_date_month	0.000000
arrival_date_week_number	0.000000

arrival_date_day_of_month	0.000000
stays_in_weekend_nights	0.000000
stays_in_week_nights	0.000000
adults	0.000000
children	0.003350
babies	0.000000
meal	0.000000
country	0.408744
market_segment	0.000000
distribution_channel	0.000000
is_repeated_guest	0.000000
previous_cancellations	0.000000
previous_bookings_not_canceled	0.000000
reserved_room_type	0.000000
assigned_room_type	0.000000
booking_changes	0.000000
deposit_type	0.000000
agent	13.686238
company	94.306893
days_in_waiting_list	0.000000
customer_type	0.000000
adr	0.000000
required_car_parking_spaces	0.000000
total_of_special_requests	0.000000
reservation_status	0.000000
reservation_status_date	0.000000
dtype:	float64

There is only one column that has more than 90% null values. This column can be dropped.

```
#Dropping column with more than 90% missing value
df = df.drop(columns = ["company"])

# Distribution of target column
plt.figure(figsize=(6,4))

ax = sns.countplot(x=df["is_canceled"])
for i in ax.containers:
    ax.bar_label(i)
ax.set(title="Cancelled vs Not cancelled")
plt.show()
```



- Out of all records, 44224 rooms have been cancelled ,i.e., 37%

In the fields present, there are children and babies. Both can be combined into a single column and the redundancy can be dropped.

```
df["kids"]=df["children"]+df["babies"]
```

```
df = df.drop(columns = ["children", "babies"])
```

```
df.describe().T
```

	count	mean	std
min \			
is_canceled	119390.0	0.370416	0.482918
0.00			
lead_time	119390.0	104.011416	106.863097
0.00			
arrival_date_year	119390.0	2016.156554	0.707476
2015.00			
arrival_date_week_number	119390.0	27.165173	13.605138
1.00			
arrival_date_day_of_month	119390.0	15.798241	8.780829
1.00			
stays_in_weekend_nights	119390.0	0.927599	0.998613
0.00			
stays_in_week_nights	119390.0	2.500302	1.908286
0.00			

adults	119390.0	1.856403	0.579261	
0.00				
is_repeated_guest	119390.0	0.031912	0.175767	
0.00				
previous_cancellations	119390.0	0.087118	0.844336	
0.00				
previous_bookings_not_canceled	119390.0	0.137097	1.497437	
0.00				
booking_changes	119390.0	0.221124	0.652306	
0.00				
agent	103050.0	86.693382	110.774548	
1.00				
days_in_waiting_list	119390.0	2.321149	17.594721	
0.00				
adr	119390.0	101.831122	50.535790	-
6.38				
required_car_parking_spaces	119390.0	0.062518	0.245291	
0.00				
total_of_special_requests	119390.0	0.571363	0.792798	
0.00				
kids	119386.0	0.111839	0.412567	
0.00				
	25%	50%	75%	max
is_canceled	0.00	0.000	1.0	1.0
lead_time	18.00	69.000	160.0	737.0
arrival_date_year	2016.00	2016.000	2017.0	2017.0
arrival_date_week_number	16.00	28.000	38.0	53.0
arrival_date_day_of_month	8.00	16.000	23.0	31.0
stays_in_weekend_nights	0.00	1.000	2.0	19.0
stays_in_week_nights	1.00	2.000	3.0	50.0
adults	2.00	2.000	2.0	55.0
is_repeated_guest	0.00	0.000	0.0	1.0
previous_cancellations	0.00	0.000	0.0	26.0
previous_bookings_not_canceled	0.00	0.000	0.0	72.0
booking_changes	0.00	0.000	0.0	21.0
agent	9.00	14.000	229.0	535.0
days_in_waiting_list	0.00	0.000	0.0	391.0
adr	69.29	94.575	126.0	5400.0
required_car_parking_spaces	0.00	0.000	0.0	8.0
total_of_special_requests	0.00	0.000	1.0	5.0
kids	0.00	0.000	0.0	10.0

In the above cell, we can see some odd values in a few columns.

Stays in weekend nights and weeknights minimum value is 0. If both are 0 in the same row, it is unrealistic. We will check for these records and drop them

```
stays = df[(df["stays_in_weekend_nights"]==0) &
(df["stays_in_week_nights"]==0)]
```

```
stays["is_canceled"].value_counts()
```

```
is_canceled
0      680
1       35
Name: count, dtype: int64
```

Here, we see that rooms were reserved and very few cancellations are present for the records but the stays are 0.

We will drop these unrealistic records

```
#Dropping records where stays in weekend nights and also weekday nights are 0
row_index=0

for i in df["stays_in_weekend_nights"]:
    if i==0 and df["stays_in_week_nights"][row_index]==0:
        df = df.drop(row_index)

        row_index=row_index+1

df = df.reset_index()
df = df.drop(columns=["index"])
```

Kids and Adults

There is also possibility of erroneous records where the bookings are done only with kids and no adults. We will filter out if there are any.

```
only_kids = df[(df["adults"]==0) & (df["kids"]!=0)]
```

```
only_kids
```

	hotel	is_canceled	lead_time	arrival_date_year	\
40598	City Hotel	0	1	2015	
40661	City Hotel	0	104	2015	
41058	City Hotel	0	3	2015	
41564	City Hotel	0	15	2015	
44762	City Hotel	1	48	2015	
...	
116493	City Hotel	0	296	2017	
116563	City Hotel	0	276	2017	
116592	City Hotel	0	291	2017	
116742	City Hotel	0	159	2017	
117487	City Hotel	0	10	2017	

	arrival_date_month	arrival_date_week_number	\
40598	August	33	
40661	August	33	

41058	August	34
41564	August	35
44762	October	43
...
116493	July	30
116563	July	31
116592	July	30
116742	July	31
117487	August	32

	arrival_date_day_of_month	stays_in_weekend_nights	\
40598	10	1	
40661	11	0	
41058	16	2	
41564	28	0	
44762	19	1	
...	
116493	27	1	
116563	30	2	
116592	29	2	
116742	31	1	
117487	12	2	

	stays_in_week_nights	adults	...	deposit_type	agent	\
40598	1	0	...	No Deposit	NaN	
40661	3	0	...	No Deposit	7.0	
41058	0	0	...	No Deposit	NaN	
41564	1	0	...	No Deposit	NaN	
44762	3	0	...	No Deposit	13.0	
...	
116493	3	0	...	No Deposit	9.0	
116563	1	0	...	No Deposit	9.0	
116592	2	0	...	No Deposit	9.0	
116742	3	0	...	No Deposit	9.0	
117487	2	0	...	No Deposit	NaN	

	days_in_waiting_list	customer_type	adr	\
40598	0	Transient-Party	9.00	
40661	0	Transient-Party	6.00	
41058	0	Transient-Party	6.00	
41564	0	Transient	0.00	
44762	0	Transient-Party	6.00	
...	
116493	0	Transient	98.85	
116563	0	Transient	93.64	
116592	0	Transient	98.85	
116742	0	Transient	121.88	
117487	0	Transient-Party	6.00	

	required_car_parking_spaces	total_of_special_requests	\
--	-----------------------------	---------------------------	---

40598	0	0
40661	0	2
41058	0	1
41564	0	1
44762	0	1
...
116493	0	1
116563	0	2
116592	0	1
116742	0	1
117487	0	1

	reservation_status	reservation_status_date	kids
40598	Check-Out	2015-08-12	3.0
40661	Check-Out	2015-08-14	2.0
41058	Check-Out	2015-08-18	2.0
41564	Check-Out	2015-08-29	2.0
44762	Canceled	2015-09-02	2.0
...
116493	Check-Out	2017-07-31	2.0
116563	Check-Out	2017-08-02	2.0
116592	Check-Out	2017-08-02	2.0
116742	Check-Out	2017-08-04	2.0
117487	Check-Out	2017-08-16	3.0

[223 rows x 30 columns]

- There are 223 rows with records where adults are 0 and kids are more than 0.
- All the records where the number of adults are 0 are unrealistic and will be dropped.

#Dropping rows where there are no adults

```
row_index=0
c=0
for i in df["adults"]:
    if i==0:
        df = df.drop(row_index)

        row_index=row_index+1

df = df.reset_index()
df = df.drop(columns=["index"])
```

Looking at values in all the columns

```
cols = df.columns
for i in cols:
    print(f"-----{i}-----")
    print(f"Values in {i} are {df[i].value_counts()}")
```

```

-----hotel-----
Values in hotel are hotel
City Hotel      78676
Resort Hotel    39666
Name: count, dtype: int64
-----is_canceled-----
Values in is_canceled are is_canceled
0      74250
1      44092
Name: count, dtype: int64
-----lead_time-----
Values in lead_time are lead_time
0      6001
1      3382
2      2042
3      1801
4      1697
...
387      1
709      1
389      1
380      1
463      1
Name: count, Length: 478, dtype: int64
-----arrival_date_year-----
Values in arrival_date_year are arrival_date_year
2016      56090
2017      40447
2015      21805
Name: count, dtype: int64
-----arrival_date_month-----
Values in arrival_date_month are arrival_date_month
August      13780
July        12553
May          11688
October     11048
April       11024
June        10879
September   10466
March        9699
February    7980
November     6706
December    6666
January     5853
Name: count, dtype: int64
-----
arrival_date_week_number-----
Values in arrival_date_week_number are arrival_date_week_number
33      3556
30      3059

```

32	3024
34	3017
18	2908
28	2821
21	2805
17	2799
20	2771
29	2750
42	2720
31	2719
41	2679
15	2668
38	2653
25	2646
27	2645
23	2605
35	2572
39	2567
22	2541
24	2486
13	2392
16	2391
40	2388
19	2377
26	2374
43	2336
14	2248
44	2243
37	2213
8	2202
36	2159
10	2131
9	2083
7	2083
11	2060
12	2055
45	1921
53	1796
49	1766
47	1669
46	1550
6	1493
48	1478
50	1473
4	1469
5	1373
3	1297
2	1201
52	1166
1	1035

```

51      909
Name: count, dtype: int64
-----
arrival_date_day_of_month-----
Values in arrival_date_day_of_month are arrival_date_day_of_month
17      4344
5       4284
15      4158
25      4133
26      4116
9       4058
12      4055
16      4044
19      4031
2       4029
20      3995
18      3980
24      3946
28      3918
8       3879
3       3830
30      3817
6       3795
14      3780
27      3759
21      3728
4       3726
13      3717
7       3624
1       3599
23      3579
22      3574
11      3559
29      3547
10      3546
31      2192
Name: count, dtype: int64
-----
stays_in_weekend_nights-----
Values in stays_in_weekend_nights are stays_in_weekend_nights
0       51160
2       33185
1       30565
4        1846
3        1251
6         152
5          77
8          58
7          19
9          10

```

10	7
12	5
13	2
16	2
18	1
19	1
14	1

Name: count, dtype: int64

-----stays_in_week_nights-----

Values in stays_in_week_nights are stays_in_week_nights

2	33611
1	30242
3	22180
5	11054
4	9516
0	6921
6	1491
10	1029
7	1024
8	654
9	226
15	85
11	55
19	43
12	42
20	39
14	35
13	27
21	15
16	15
22	7
18	6
25	6
17	4
30	4
24	3
40	2
33	1
50	1
42	1
32	1
26	1
34	1

Name: count, dtype: int64

-----adults-----

Values in adults are adults

2	89253
1	22825
3	6186
4	62

```

26      5
5       2
27      2
20      2
40      1
55      1
50      1
6       1
10      1
Name: count, dtype: int64
-----meal-----
Values in meal are meal
BB      91505
HB      14380
SC      10500
Undefined 1160
FB       797
Name: count, dtype: int64
-----country-----
Values in country are country
PRT     47900
GBR     12080
FRA     10364
ESP      8530
DEU      7271
...
NAM       1
GUY       1
LCA       1
MRT       1
ASM       1
Name: count, Length: 177, dtype: int64
-----market_segment-----
Values in market_segment are market_segment
Online TA      55942
Offline TA/T0  24046
Groups         19759
Direct         12420
Corporate       5231
Complementary   711
Aviation        231
Undefined        2
Name: count, dtype: int64
-----distribution_channel-----
Values in distribution_channel are distribution_channel
TA/T0      97145
Direct     14418
Corporate   6584
GDS         190
Undefined    5

```

```

Name: count, dtype: int64
-----is_repeated_guest-----
Values in is_repeated_guest are is_repeated_guest
0      114844
1       3498
Name: count, dtype: int64
-----
previous_cancellations-----
Values in previous_cancellations are previous_cancellations
0      111869
1       6042
2        114
3         65
24         48
11         35
4          31
26         26
25         25
6          22
5          19
19         19
14         14
13         12
21          1
Name: count, dtype: int64
-----
previous_bookings_not_canceled-----
Values in previous_bookings_not_canceled are
previous_bookings_not_canceled
0      114760
1       1519
2        576
3        331
4        226
...
60         1
61         1
62         1
63         1
72         1
Name: count, Length: 73, dtype: int64
-----reserved_room_type-----
Values in reserved_room_type are reserved_room_type
A      85388
D      19095
E       6481
F       2877
G       2073
C        923
B        902

```

```

H      597
L        6
Name: count, dtype: int64
-----assigned_room_type-----
Values in assigned_room_type are assigned_room_type
A      73752
D      25202
E       7763
F       3728
G       2531
C       2350
B       1961
H        708
I        219
K        127
L         1
Name: count, dtype: int64
-----booking_changes-----
Values in booking_changes are booking_changes
0      100548
1      12529
2       3752
3        906
4        364
5        110
6         60
7         27
8         14
9          8
10         6
13         5
15         3
14         3
17         2
16         2
12         1
18         1
11         1
Name: count, dtype: int64
-----deposit_type-----
Values in deposit_type are deposit_type
No Deposit    103593
Non Refund    14587
Refundable     162
Name: count, dtype: int64
-----agent-----
Values in agent are agent
9.0      31676
240.0    13795
1.0      7185

```



```

14.0      3604
7.0       3515
...
59.0       1
451.0      1
358.0      1
397.0      1
433.0      1
Name: count, Length: 333, dtype: int64
-----days_in_waiting_list-----
Values in days_in_waiting_list are days_in_waiting_list
0      114650
39       227
58       164
44       141
31       127
...
109       1
37        1
193       1
73        1
36        1
Name: count, Length: 127, dtype: int64
-----customer_type-----
Values in customer_type are customer_type
Transient      88796
Transient-Party 24923
Contract       4055
Group          568
Name: count, dtype: int64
-----adr-----
Values in adr are adr
62.00      3754
75.00      2714
90.00      2472
65.00      2418
80.00      1889
...
199.77      1
213.81      1
90.70       1
95.12       1
173.73      1
Name: count, Length: 8859, dtype: int64
-----
required_car_parking_spaces-----
Values in required_car_parking_spaces are required_car_parking_spaces
0      110951
1       7358
2         28

```

```

3      3
8      2
Name: count, dtype: int64
-----
total_of_special_requests-----
Values in total_of_special_requests are total_of_special_requests
0      69732
1      32907
2      12853
3       2472
4        338
5         40
Name: count, dtype: int64
-----reservation_status-----
Values in reservation_status are reservation_status
Check-Out      74250
Canceled       42902
No-Show        1190
Name: count, dtype: int64
-----
reservation_status_date-----
Values in reservation_status_date are reservation_status_date
2015-10-21      1456
2015-07-06       803
2016-11-25       787
2015-01-01       763
2016-01-18       623
...
2015-04-21        1
2015-03-11        1
2015-03-12        1
2015-03-18        1
2017-09-12        1
Name: count, Length: 926, dtype: int64
-----kids-----
Values in kids are kids
0.0      109268
1.0       5418
2.0       3552
3.0         97
10.0         2
9.0          1
Name: count, dtype: int64

```

- Lead time, stays_in_weeknights, stays_in_weekend, waiting list and adr are the only three numerical columns
- Although many other columns have numerical values, they are categorical in nature

```

#Diving columns into numerical ad categorical
num_cols = ["lead_time", 'stays_in_weekend_nights',
            'stays_in_week_nights', 'days_in_waiting_list', 'adr']
cat_cols = ['hotel', 'is_canceled',
            'arrival_date_year', 'arrival_date_month',
            'arrival_date_day_of_month', 'adults', 'meal',
            'country', 'market_segment',
            'distribution_channel', 'is_repeated_guest', 'previous_cancellations',
            'previous_bookings_not_canceled',
            'reserved_room_type', 'assigned_room_type', 'booking_changes',
            'deposit_type', 'agent',
            'customer_type', 'required_car_parking_spaces',
            'total_of_special_requests', 'reservation_status',
            'reservation_status_date', 'kids']

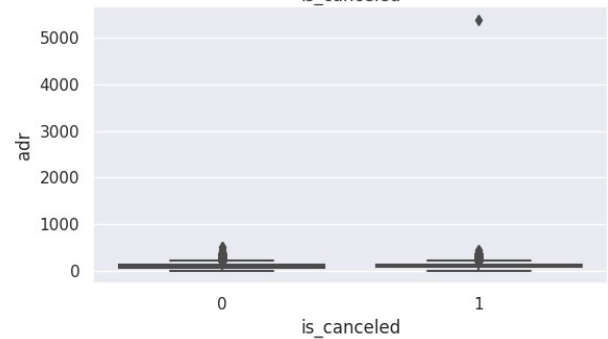
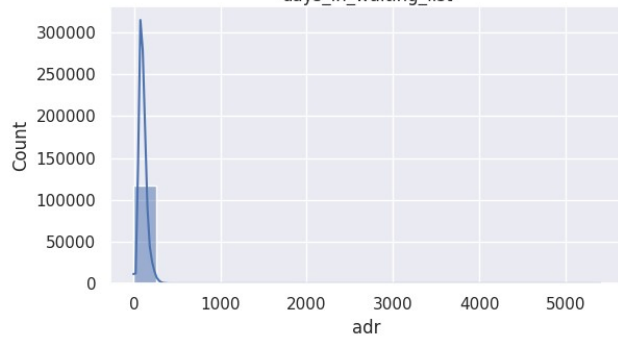
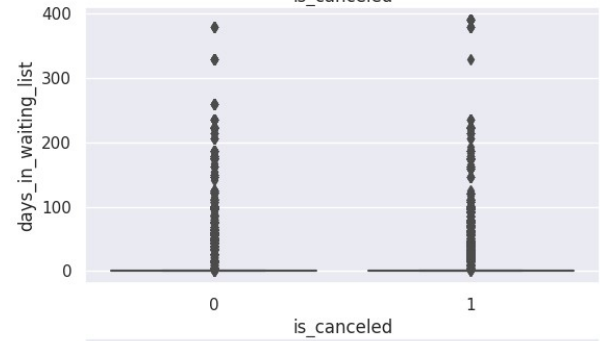
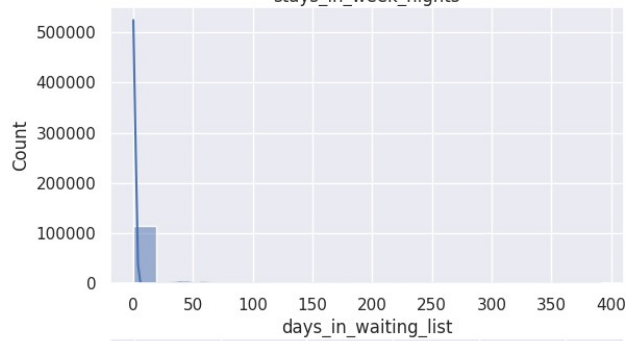
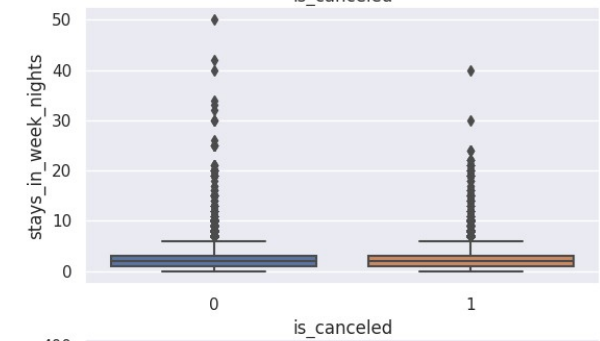
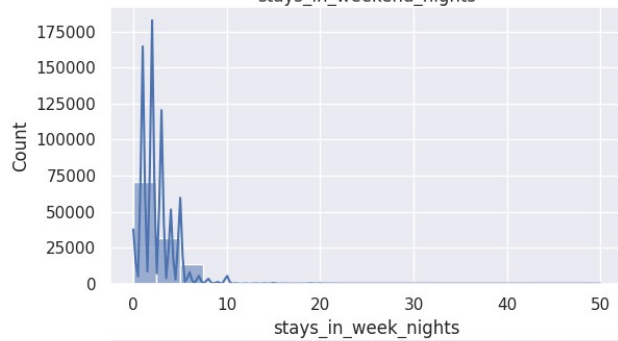
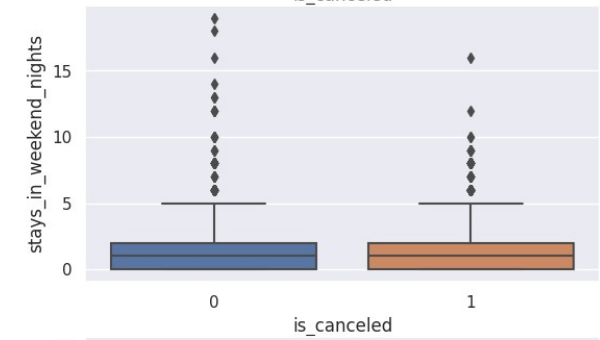
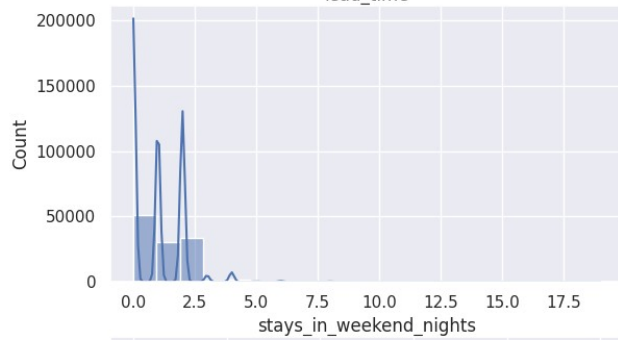
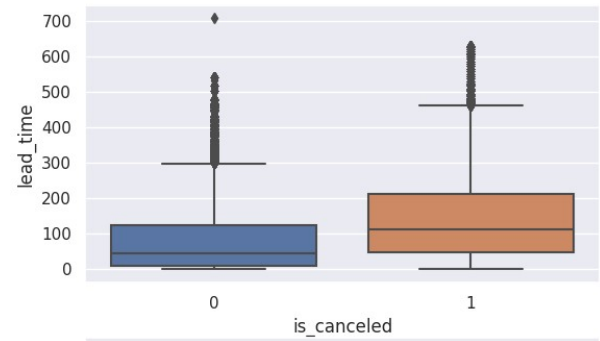
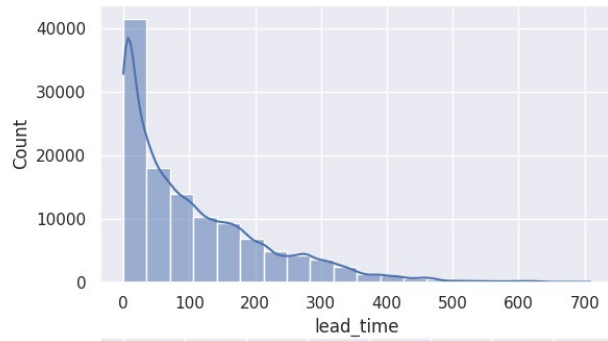
```

Numerical columns

```

#plotting distribution of numerical columns
r = 5
c = 2
f, axes = plt.subplots(r, c, figsize=(14, 20))
ctr = 0
for i in range(r):
    #for j in range(c):
        j=0
        col = num_cols[ctr]
        sns.histplot(x=df[col], bins=20, kde=True, ax=axes[i, j])
        sns.boxplot(x=df["is_canceled"], y=df[col], ax=axes[i, j+1])
        ctr = ctr + 1
plt.show()

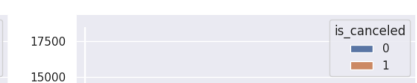
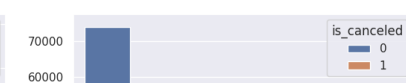
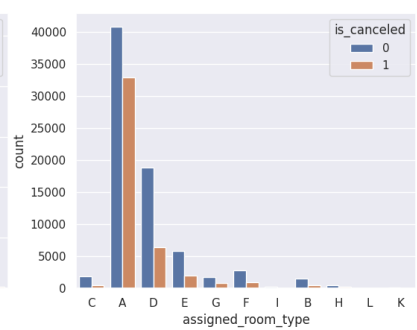
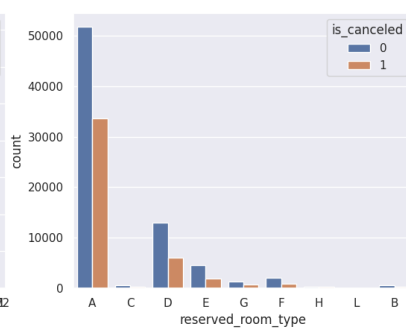
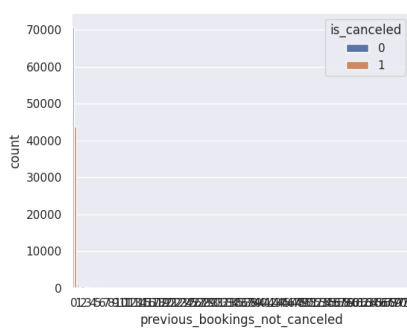
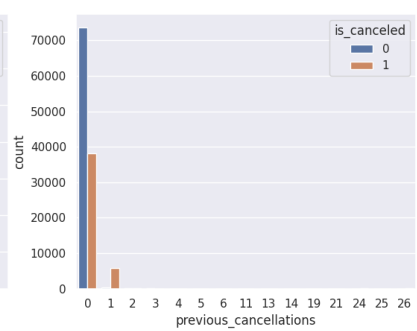
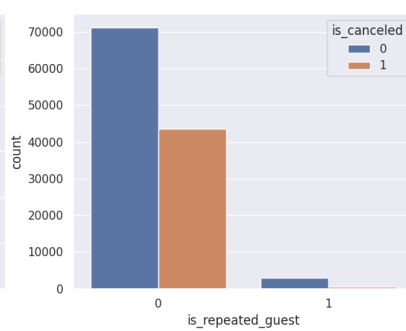
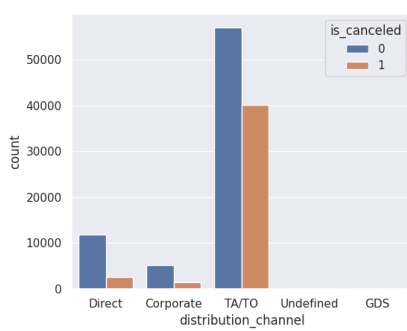
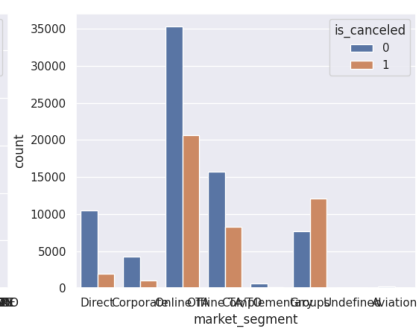
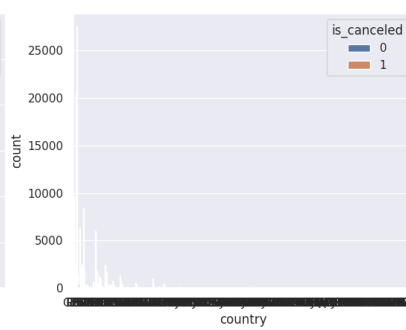
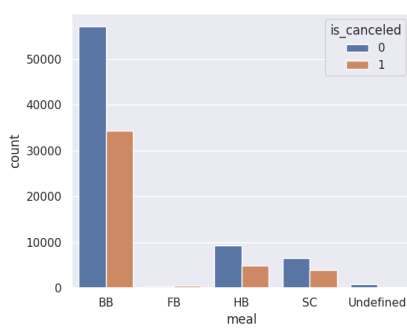
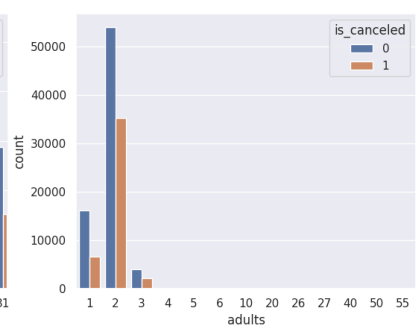
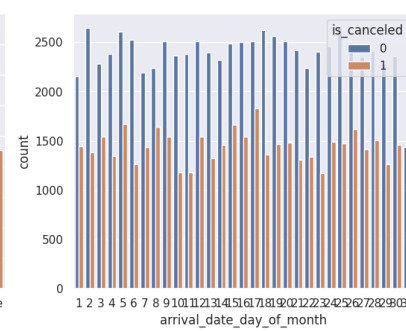
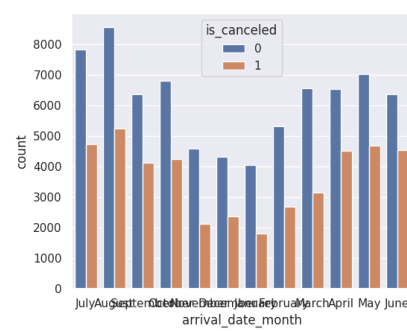
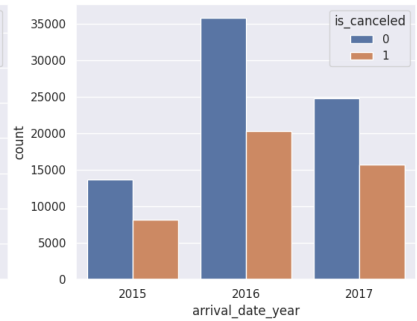
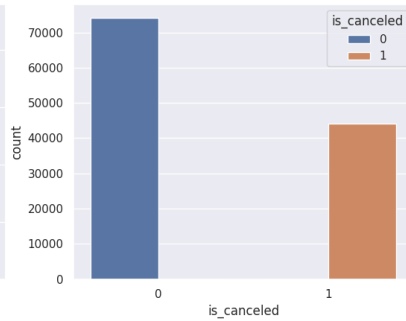
```



- There are significant number of outliers in adr, days in waiting list, stays in weeknights and weekend nights.
- Outliers in waiting list column and adr need to be addressed as they are very high and may significantly impact the outcome.

Categorical columns

```
#Plotting distribution of categorical columns
r = 8
c = 3
f, axes = plt.subplots(r, c, figsize=(20, 45))
ctr = 0
for i in range(r):
    for j in range(c):
        col = cat_cols[ctr]
        sns.countplot(x=df[col], hue=df["is_canceled"], ax=axes[i, j])
        ctr = ctr + 1
plt.show()
```



- Bookings and cancellations in a city hotel is higher than resort hotel
- Bookings in year 2016 is highest because of the data distribution. We have data from July 2015 to June 2017.
- There is a significant dip in bookings in November, December and January. As summer in most inhabited continents of the world lasts from May to September, we can assume that for a tourist destination like Portugal, November to January will be a low season and May to September will be high season.
- There is no discernable pattern between day of the month and cancellation.
- The maximum number of bookings seem to be amongst 1 to 3 adults where 2 adults is the highest.
- There are different types of meals available for the occupants. However, it can be assumed that this will not have impact on the cancellations as there are other more significant factors involved.
- There are total of 177 countries from which tourists are visiting. This needs to be examined further.
- market_segment and distribution_channel are interrelated columns. We can keep market_segment for further exploration.
- is_repeated_guest, previous_cancellations and previous_bookings_not_cancelled data is highly skewed. Needs further exploration before being dropped.
- room A has highest bookings. We must check if cancellations are high where reserved and assigned room types are not same.
- booking_changes also has highly skewed data.
- Cancellations are high for non-refundable deposit types.
- There is no particular description for agent column. It appears to be a number associated with the agent responsible for the booking. This column can be dropped.
- reservation_status and reservation_status_date can be dropped as they are redundant.

Checking data and cancellation percentages for columns with skewed data.

```
perc = df.groupby(['is_repeated_guest',
'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
perc
```

is_canceled	0	1	total perc	canc perc
is_repeated_guest				
0	71301	43543	0.970442	0.379149
1	2949	549	0.029558	0.156947

More than 96% of the data is for one value only. We will drop this column as it is skewed.

```
perc = df.groupby(['previous_cancellations',
'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
perc
```

is_canceled	0	1	total perc	canc perc
previous_cancellations				
0	73715.0	38154.0	0.945303	0.341060
1	332.0	5710.0	0.051055	0.945051
2	76.0	38.0	0.000963	0.333333
3	45.0	20.0	0.000549	0.307692
4	24.0	7.0	0.000262	0.225806
5	17.0	2.0	0.000161	0.105263
6	15.0	7.0	0.000186	0.318182
11	25.0	10.0	0.000296	0.285714
13	1.0	11.0	0.000101	0.916667
14	NaN	14.0	NaN	NaN
19	NaN	19.0	NaN	NaN
21	NaN	1.0	NaN	NaN
24	NaN	48.0	NaN	NaN
25	NaN	25.0	NaN	NaN
26	NaN	26.0	NaN	NaN

```

perc = df.groupby(['previous_bookings_not_canceled',
'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
perc

```

is_canceled	0	1	total perc	canc
perc				
previous_bookings_not_canceled				
0	70867.0	43893.0	0.969732	
0.382476				
1	1441.0	78.0	0.012836	
0.051350				
2	544.0	32.0	0.004867	
0.055556				
3	314.0	17.0	0.002797	
0.051360				
4	214.0	12.0	0.001910	
0.053097				
...
.				
68	1.0	NaN	NaN	
NaN				
69	1.0	NaN	NaN	
NaN				
70	1.0	NaN	NaN	
NaN				
71	1.0	NaN	NaN	
NaN				
72	1.0	NaN	NaN	
NaN				

[73 rows x 4 columns]

More than 96% of the data is for one value only. We will drop this column as it is skewed.

```
perc = df.groupby(['booking_changes',  
                  'is_canceled']).size().unstack(1)  
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])  
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))  
perc
```

is_canceled	0	1	total perc	canc perc
booking_changes				
0	59268.0	41280.0	0.849639	0.410550
1	10737.0	1792.0	0.105871	0.143028
2	2991.0	761.0	0.031705	0.202825
3	763.0	143.0	0.007656	0.157837
4	297.0	67.0	0.003076	0.184066
5	90.0	20.0	0.000930	0.181818
6	42.0	18.0	0.000507	0.300000
7	24.0	3.0	0.000228	0.111111
8	10.0	4.0	0.000118	0.285714
9	7.0	1.0	0.000068	0.125000
10	5.0	1.0	0.000051	0.166667
11	1.0	NaN	NaN	NaN
12	1.0	NaN	NaN	NaN
13	5.0	NaN	NaN	NaN
14	2.0	1.0	0.000025	0.333333
15	3.0	NaN	NaN	NaN
16	1.0	1.0	0.000017	0.500000
17	2.0	NaN	NaN	NaN
18	1.0	NaN	NaN	NaN

```
perc = df.groupby(['required_car_parking_spaces',  
                  'is_canceled']).size().unstack(1)  
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])  
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))  
perc
```

is_canceled	0	1	total perc	canc perc
required_car_parking_spaces				
0	66859.0	44092.0	0.937545	0.397401
1	7358.0	NaN	NaN	NaN
2	28.0	NaN	NaN	NaN
3	3.0	NaN	NaN	NaN
8	2.0	NaN	NaN	NaN

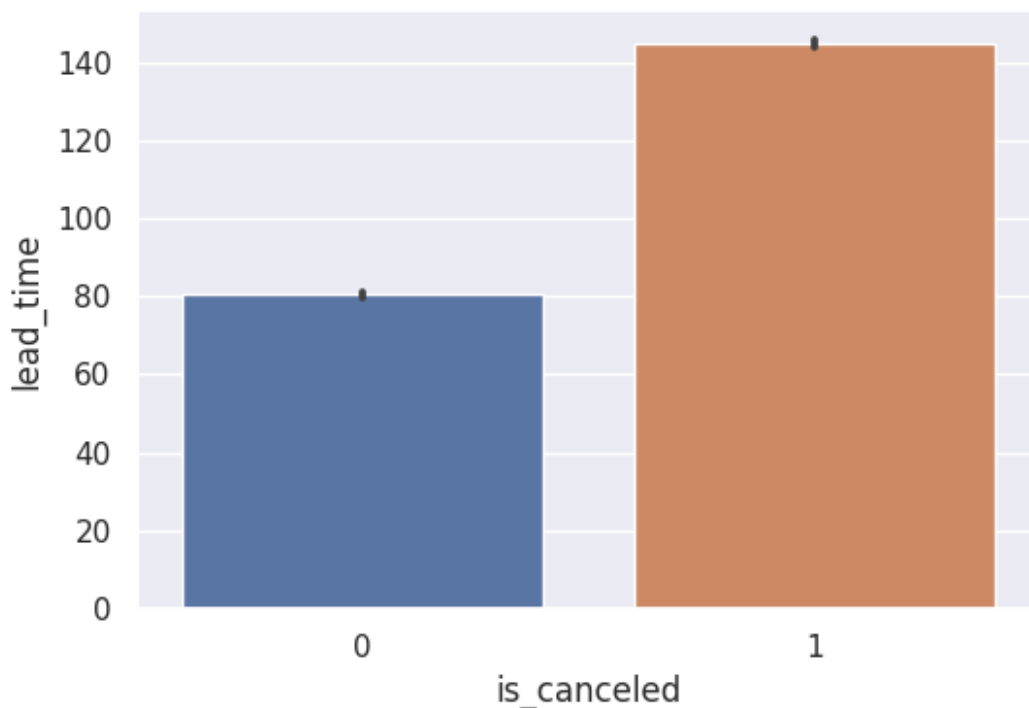
```
#Dropping columns with lesser impact  
cols_to_drop = ['reservation_status',
```

```
'reservation_status_date', 'agent', 'distribution_channel', 'meal',
    'previous_bookings_not_canceled', 'is_repeated_guest']
df.drop(columns = cols_to_drop, inplace=True)
```

Lead Time and waiting list

We may assume that if there is higher lead time, there is a higher chance of cancellation as people may have longer time to change options. We will check this.

```
sns.set(rc={"figure.figsize":(6,4)})
sns.barplot(x=df["is_canceled"],y=df["lead_time"])
<Axes: xlabel='is_canceled', ylabel='lead_time'>
```



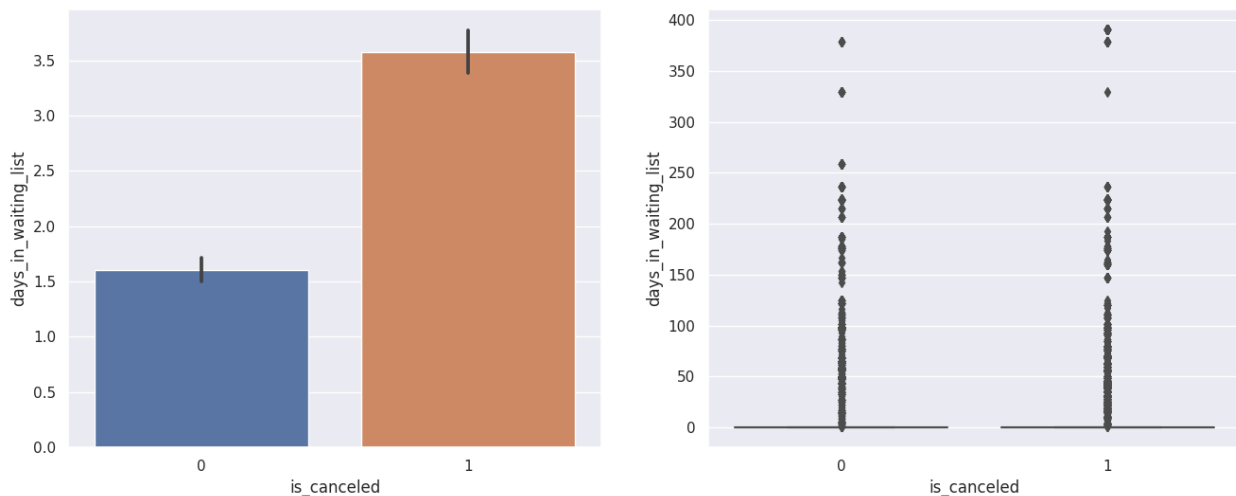
Our assumption was correct. Higher lead time is leading to higher cancellations.

Similarly, we may also assume that higher number of days in waiting list may also lead to higher cancellations.

```
sns.set(rc={"figure.figsize":(16,6)})
f, axes = plt.subplots(1, 2)

sns.barplot(x=df["is_canceled"],y=df["days_in_waiting_list"],ax=axes[0])
sns.boxplot(x=df["is_canceled"],y=df["days_in_waiting_list"],ax=axes[1])
```

```
plt.show()
```



- This assumption is also true. Higher days in waiting list lead to higher cancellations.
- The mean number of days in waiting list is around 3.5 but there are outliers that go up to 400 days in waiting list.

We can drop week number column as year, month and day is already given

```
df = df.drop(columns = ["arrival_date_week_number"])
```

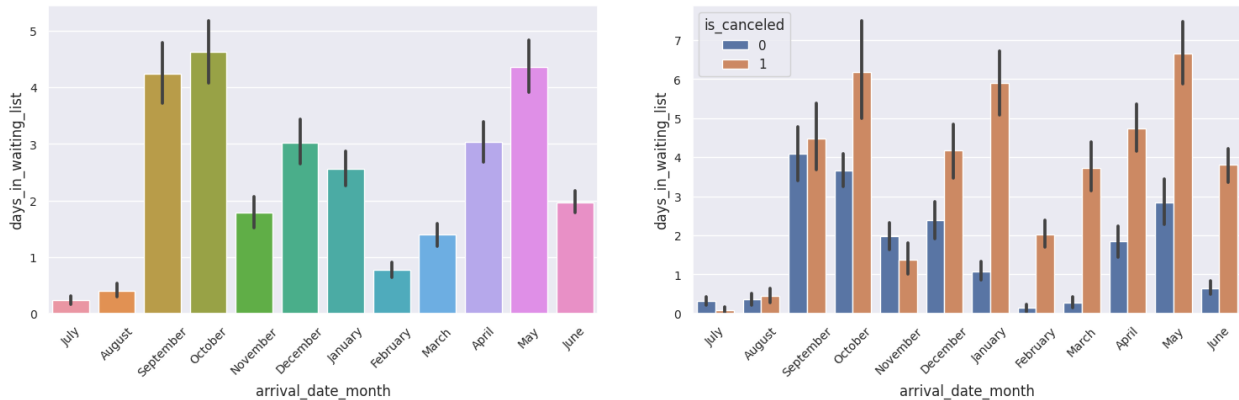
Relationship between month and days in waiting list.

```
# Set figure size
sns.set(rc={"figure.figsize":(16,6)})
f, axes = plt.subplots(1, 2, figsize=(16, 6))

# Plotting barplots
sns.barplot(x=df["arrival_date_month"], y=df["days_in_waiting_list"],
ax=axes[0])
sns.barplot(x=df["arrival_date_month"], y=df["days_in_waiting_list"],
hue=df["is_canceled"], ax=axes[1])

# Adjust font size and rotation of labels
for ax in axes:
    ax.tick_params(axis='both', which='major', labelsz=10) # Reduce
font size of tick labels
    ax.set_xlabel(ax.get_xlabel(), fontsize=12) # Adjust font size of
x-axis labels
    ax.set_ylabel(ax.get_ylabel(), fontsize=12) # Adjust font size of
y-axis labels
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45) # Rotate x-
axis labels to prevent overlap
```

```
plt.tight_layout(pad=4.0) # Adjust spacing between subplots
plt.show()
```



- The peak tourist season as per many travel websites is between July and September.
- We can assume that since October is the end of tourist season and May being the beginning of tourist season, the mean number of days in waiting list is high.
- We can also see that despite lesser mean number of days in waiting list in January, February, March and June, the mean days in waiting list for cancellation is very high. This can indicate that if days in waiting list crosses a certain threshold especially during off-season, there is higher tendency of cancellation.

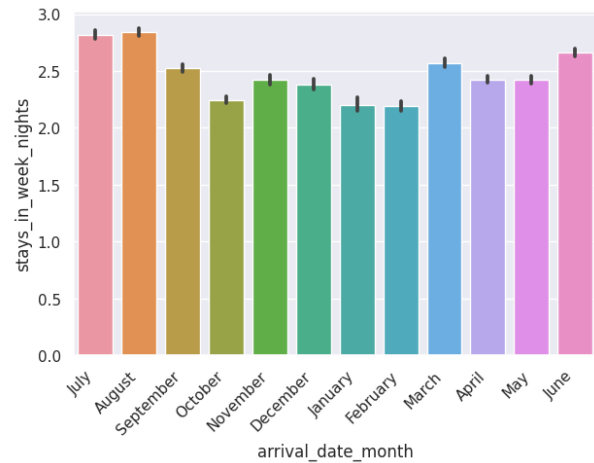
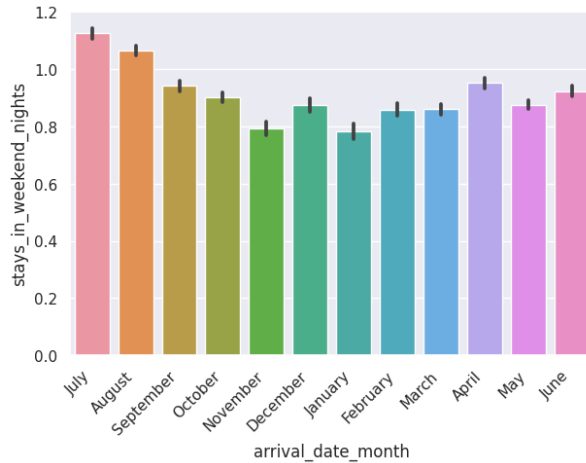
Duration and type of stay

```
# Set a more appropriate figure size
sns.set(rc={"figure.figsize":(14,6)}) # Adjust width and height to
better fit the plots
f, axes = plt.subplots(1, 2, figsize=(14, 6)) # Adjust figure size

# Plotting barplots
sns.barplot(x=df["arrival_date_month"],
y=df["stays_in_weekend_nights"], ax=axes[0])
sns.barplot(x=df["arrival_date_month"], y=df["stays_in_week_nights"],
ax=axes[1])

# Rotate x-axis labels to prevent overlap
for ax in axes:
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
# Rotate labels and align them to the right

plt.tight_layout(pad=4.0) # Adjust spacing between subplots
plt.show()
```



The average number of stays in weeknights is more than weekend nights. This is probably because there are more number of weeknights. However, the overall pattern of bookings look similar for both.

We want to see if there is any relation between cancellation and duration of stay during weekends or weeknights.

```
df["total_stay"] = df["stays_in_weekend_nights"] +
df["stays_in_week_nights"]

#Dividing data into different columns based on the duration of stay
df["weekend"] = 0
df["weekday"] = 0
df["full_week"] = 0
df["long_weekend"] = 0
df["long_stay"] = 0

row_index = 0

for i in df["stays_in_weekend_nights"]:
    if df["total_stay"][row_index] < 7:
        if i == 0 and df["stays_in_week_nights"][row_index] > 0:
            df["weekday"][row_index] = 1
        if i > 0 and df["stays_in_week_nights"][row_index] == 0:
            df["weekend"][row_index] = 1
        if i > 0 and df["stays_in_week_nights"][row_index] > 0:
            df["long_weekend"][row_index] = 1
        if i > 0 and df["stays_in_week_nights"][row_index] > 0 and
df["total_stay"][row_index] == 7:
            df["full_week"][row_index] = 1

        if df["total_stay"][row_index] >= 7:
            df["long_stay"][row_index] = 1
        row_index = row_index + 1
```

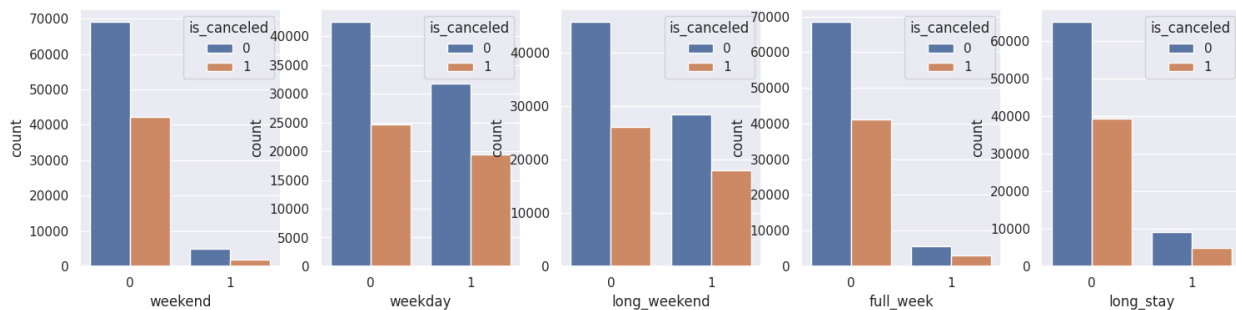
```

sns.set(rc={"figure.figsize":(18,4)})
f, axes = plt.subplots(1,5)

sns.countplot(x=df["weekend"],hue = df["is_canceled"],ax=axes[0])
sns.countplot(x=df["weekday"],hue = df["is_canceled"],ax=axes[1])
sns.countplot(x=df["long_weekend"],hue = df["is_canceled"],ax=axes[2])
sns.countplot(x=df["full_week"],hue = df["is_canceled"],ax=axes[3])
sns.countplot(x=df["long_stay"],hue = df["is_canceled"],ax=axes[4])

plt.show()

```



```

stay_dur =
['weekend','weekday','long_weekend','full_week','long_stay']
for stay in stay_dur:
    perc = df.groupby([stay, 'is_canceled']).size().unstack(1)
    perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
    perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
    print(perc)

```

is_canceled	0	1	total perc	canc perc
weekend				
0	69205	42216	0.941517	0.378887
1	5045	1876	0.058483	0.271059
is_canceled	0	1	total perc	canc perc
weekday				
0	42521	24661	0.567694	0.367077
1	31729	19431	0.432306	0.379808
is_canceled	0	1	total perc	canc perc
long_weekend				
0	45795	26145	0.607899	0.363428
1	28455	17947	0.392101	0.386772
is_canceled	0	1	total perc	canc perc
full_week				
0	68574	41129	0.927	0.374912
1	5676	2963	0.073	0.342980
is_canceled	0	1	total perc	canc perc
long_stay				
0	65229	39254	0.88289	0.375697
1	9021	4838	0.11711	0.349087

- We can see that a stay of full week or more has lesser bookings but significant cancellations.
- Only weekend stay is also less common and cancellations are also less.
- But staying on weekdays or weekdays including weekends is more common and the cancellations percentages are also significant.
- We can drop the redundant columns

```
#Dropping redundant columns related to stay
df = df.drop(columns =
["stays_in_weekend_nights", "stays_in_week_nights", 'arrival_date_day_of_month'])
```

Country

```
cnt = dict(df['country'].value_counts(normalize=True))
cnt

{'PRT': 0.40638335779551876,
 'GBR': 0.10248665891795128,
 'FRA': 0.08792812359483834,
 'ESP': 0.0723684768683878,
 'DEU': 0.061687127234472165,
 'ITA': 0.031789529053440686,
 'IRL': 0.02858257896478294,
 'BEL': 0.01969983625889759,
 'BRA': 0.018724176840390603,
 'NLD': 0.017748517421883617,
 'USA': 0.017655193477504688,
 'CHE': 0.014550051328169408,
 'CN': 0.010842545537842859,
 'AUT': 0.010681349633915618,
 'SWE': 0.008619738862635638,
 'CHN': 0.008458542958708397,
 'POL': 0.007737403388507581,
 'ISR': 0.005667308622284061,
 'RUS': 0.005285528849824805,
 'NOR': 0.005132816940841103,
 'ROU': 0.00421654548693889,
 'FIN': 0.003766893754931322,
 'DNK': 0.003648117825721776,
 'AUS': 0.00361418184594762,
 'AGO': 0.0030627221746175838,
 'LUX': 0.00240945456396508,
 'MAR': 0.0021634187106024487,
 'TUR': 0.002095546751054136,
 'HUN': 0.0019343508471268952,
 'ARG': 0.0017986069280302708,
```

'JPN': 0.0016713470038771858,
'CZE': 0.0014507631353451714,
'IND': 0.001272599241530852,
'KOR': 0.0011198873325471497,
'GRC': 0.0010859513527729937,
'DZA': 0.0008738514791845184,
'SRB': 0.0008568834892974404,
'HRV': 0.0008483994943539014,
'EST': 0.0007041715803137381,
'MEX': 0.0007041715803137381,
'IRN': 0.0006956875853701991,
'LTU': 0.0006872035904266601,
'ZAF': 0.0006702356005395821,
'BGR': 0.000636299620765426,
'NZL': 0.000627815625821887,
'COL': 0.00060236364099127,
'UKR': 0.0005769116561606529,
'MOZ': 0.0005599436662735749,
'SVK': 0.0005514596713300359,
'CHL': 0.0005429756763864968,
'THA': 0.0004920717067252628,
'ISL': 0.0004751037168381848,
'LVA': 0.00046661972189464577,
'SVN': 0.00046661972189464577,
'TWN': 0.0004326837421204897,
'CYP': 0.0004326837421204897,
'ARE': 0.00041571575223341167,
'SAU': 0.00039874776234633365,
'PHL': 0.0003935979774156054,
'SGP': 0.0003223918078544825,
'TUN': 0.0003223918078544825,
'IDN': 0.0002969398230238655,
'NGA': 0.00028845582808032644,
'URY': 0.0002714878381932484,
'EGY': 0.0002630038432497094,
'LBN': 0.0002630038432497094,
'HKG': 0.0002460358533626314,
'PER': 0.0002460358533626314,
'MYS': 0.0002375518584190924,
'ECU': 0.00022906786347555338,
'VEN': 0.00022058386853201435,
'BLR': 0.00020361587864493633,
'CPV': 0.00020361587864493633,
'GEO': 0.0001866478887578583,
'JOR': 0.00016967989887078027,
'KAZ': 0.00016119590392724126,
'CRI': 0.00016119590392724126,
'GIB': 0.00015271190898370225,
'MLT': 0.00015271190898370225,

'AZE' : 0.00014422791404016322,
'OMN' : 0.00014422791404016322,
'MAC' : 0.0001357439190966242,
'KWT' : 0.0001357439190966242,
'QAT' : 0.0001187759292095462,
'DOM' : 0.0001187759292095462,
'IRQ' : 0.0001187759292095462,
'PAK' : 0.0001187759292095462,
'BIH' : 0.00011029193426600717,
'PRI' : 0.00010180793932246817,
'BGD' : 0.00010180793932246817,
'MDV' : 0.00010180793932246817,
'ALB' : 0.00010180793932246817,
'SEN' : 9.332394437892914e-05,
'BOL' : 8.483994943539014e-05,
'MKD' : 8.483994943539014e-05,
'CMR' : 8.483994943539014e-05,
'TJK' : 7.635595449185113e-05,
'GNB' : 7.635595449185113e-05,
'PAN' : 7.635595449185113e-05,
'ARM' : 6.78719595483121e-05,
'CUB' : 6.78719595483121e-05,
'VNM' : 6.78719595483121e-05,
'JEY' : 6.78719595483121e-05,
'LBY' : 6.78719595483121e-05,
'AND' : 5.93879646047731e-05,
'LKA' : 5.93879646047731e-05,
'MUS' : 5.93879646047731e-05,
'CIV' : 5.090396966123408e-05,
'JAM' : 5.090396966123408e-05,
'KEN' : 5.090396966123408e-05,
'MNE' : 4.241997471769507e-05,
'SUR' : 4.241997471769507e-05,
'FRO' : 4.241997471769507e-05,
'BHR' : 4.241997471769507e-05,
'CAF' : 4.241997471769507e-05,
'GTM' : 3.393597977415605e-05,
'GAB' : 3.393597977415605e-05,
'ZWE' : 3.393597977415605e-05,
'MCO' : 3.393597977415605e-05,
'GHA' : 3.393597977415605e-05,
'UZB' : 3.393597977415605e-05,
'PRY' : 3.393597977415605e-05,
'BRB' : 3.393597977415605e-05,
'BEN' : 2.545198483061704e-05,
'TMP' : 2.545198483061704e-05,
'ETH' : 2.545198483061704e-05,
'TZA' : 2.545198483061704e-05,
'SYR' : 2.545198483061704e-05,

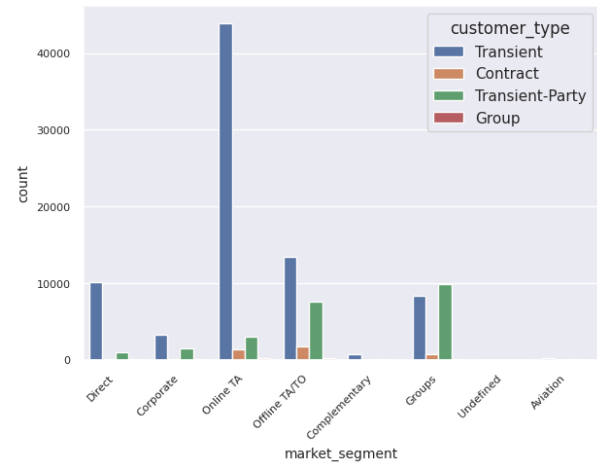
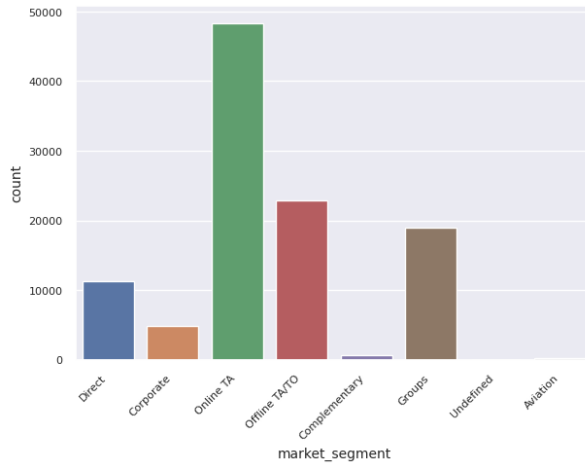
'GGY' : 2.545198483061704e-05,
'LIE' : 2.545198483061704e-05,
'LAO' : 1.6967989887078026e-05,
'SLV' : 1.6967989887078026e-05,
'ABW' : 1.6967989887078026e-05,
'MYT' : 1.6967989887078026e-05,
'ATA' : 1.6967989887078026e-05,
'GLP' : 1.6967989887078026e-05,
'MWI' : 1.6967989887078026e-05,
'UGA' : 1.6967989887078026e-05,
'COM' : 1.6967989887078026e-05,
'TGO' : 1.6967989887078026e-05,
'ZMB' : 1.6967989887078026e-05,
'KNA' : 1.6967989887078026e-05,
'RWA' : 1.6967989887078026e-05,
'STP' : 1.6967989887078026e-05,
'SYC' : 1.6967989887078026e-05,
'KHM' : 1.6967989887078026e-05,
'IMN' : 1.6967989887078026e-05,
'NCL' : 8.483994943539013e-06,
'KIR' : 8.483994943539013e-06,
'SDN' : 8.483994943539013e-06,
'ATF' : 8.483994943539013e-06,
'SLE' : 8.483994943539013e-06,
'AIA' : 8.483994943539013e-06,
'DMA' : 8.483994943539013e-06,
'PYF' : 8.483994943539013e-06,
'NIC' : 8.483994943539013e-06,
'VGB' : 8.483994943539013e-06,
'HND' : 8.483994943539013e-06,
'NPL' : 8.483994943539013e-06,
'BHS' : 8.483994943539013e-06,
'DJI' : 8.483994943539013e-06,
'PLW' : 8.483994943539013e-06,
'BWA' : 8.483994943539013e-06,
'MDG' : 8.483994943539013e-06,
'UMI' : 8.483994943539013e-06,
'SMR' : 8.483994943539013e-06,
'BDI' : 8.483994943539013e-06,
'FJI' : 8.483994943539013e-06,
'CYM' : 8.483994943539013e-06,
'MMR' : 8.483994943539013e-06,
'BFA' : 8.483994943539013e-06,
'MLI' : 8.483994943539013e-06,
'NAM' : 8.483994943539013e-06,
'GUY' : 8.483994943539013e-06,
'LCA' : 8.483994943539013e-06,
'MRT' : 8.483994943539013e-06,
'ASM' : 8.483994943539013e-06}

We can see that top 15 countries cover more than 80% of the data. For ease of exploration and processing, we will only keep the top 15 countries.

```
df = df[(df["country"]=="PRT") | (df["country"]=="GBR") |  
(df["country"]=="FRA") | (df["country"]=="ESP") |  
(df["country"]=="DEU") |  
        (df["country"]=="ITA") | (df["country"]=="IRL") |  
(df["country"]=="BEL") | (df["country"]=="BRA") |  
        (df["country"]=="NLD") | (df["country"]=="USA") |  
(df["country"]=="CHE") | (df["country"]=="CN") |  
        (df["country"]=="AUT") | (df["country"]=="SWE")]
```

Market Segment and Customer type

```
# Dropping undefined values from market segment  
row_index = 0  
for i in df["market_segment"]:  
    if i == "Undefined":  
        df = df.drop(row_index)  
        row_index = row_index + 1  
  
df = df.reset_index()  
df = df.drop(columns=["index"])  
  
# Set figure size and create subplots  
sns.set(rc={"figure.figsize":(14,6)}) # Adjust figure size as needed  
f, axes = plt.subplots(1, 2, figsize=(14, 6))  
  
# Plotting countplots  
sns.countplot(x=df["market_segment"], ax=axes[0])  
sns.countplot(x=df["market_segment"], hue=df["customer_type"],  
ax=axes[1])  
  
# Adjust font size of ticks and labels  
for ax in axes:  
    ax.tick_params(axis='both', which='major', labelsize=8) # Reduce  
font size of tick labels  
    ax.set_xlabel(ax.get_xlabel(), fontsize=10) # Reduce font size of  
x-axis labels  
    ax.set_ylabel(ax.get_ylabel(), fontsize=10) # Reduce font size of  
y-axis labels  
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')  
# Rotate x-axis labels  
  
plt.tight_layout(pad=4.0) # Adjust spacing between subplots  
plt.show()
```



```
perc = df.groupby(['market_segment', 'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
print(perc)
```

is_canceled	0	1	total perc	canc perc
market_segment				
Aviation	152.0	50.0	0.001884	0.247525
Complementary	609.0	85.0	0.006472	0.122478
Corporate	3871.0	972.0	0.045165	0.200702
Direct	9518.0	1739.0	0.104981	0.154482
Groups	6886.0	12056.0	0.176650	0.636469
Offline TA/TO	14656.0	8267.0	0.213776	0.360642
Online TA	30831.0	17535.0	0.451053	0.362548
Undefined	NaN	2.0	NaN	NaN

- Under market segment, the highest booking is by Online Travel Agents. However the highest cancellation percentage is with the Groups.

```
# Adjusting figure size to prevent scrolling
sns.set(rc={"figure.figsize":(12,6)}) # Smaller figure size
f, axes = plt.subplots(1, 2, figsize=(12, 6)) # Adjust size for two
plots side by side

# Plotting countplots
sns.countplot(x=df["market_segment"], hue=df["is_canceled"],
ax=axes[0])
sns.countplot(x=df["customer_type"], hue=df["is_canceled"],
ax=axes[1])

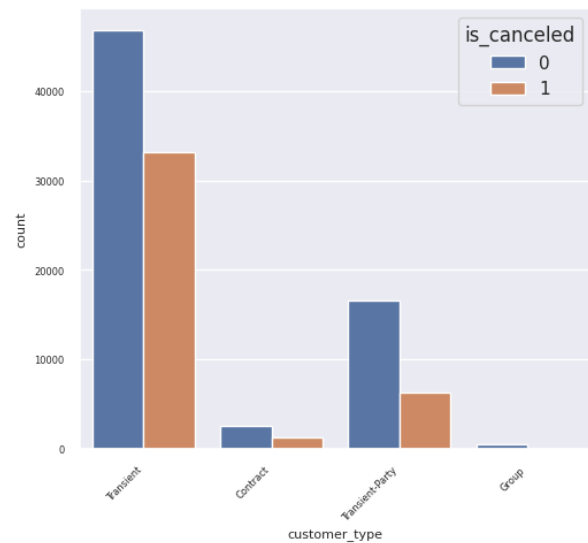
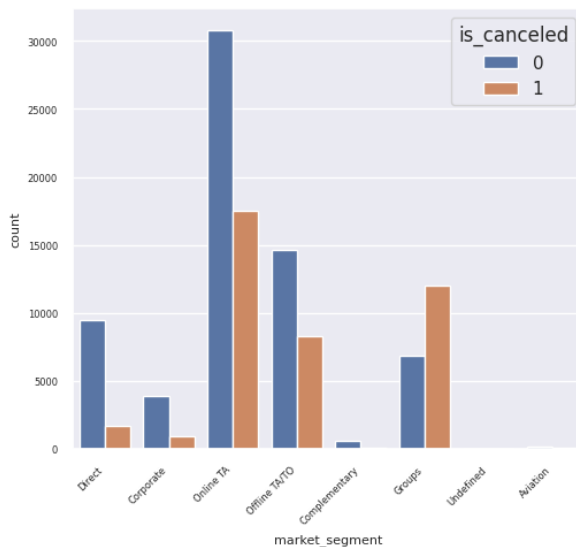
# Adjust font size and rotate labels
for ax in axes:
    ax.tick_params(axis='both', which='major', labelsz=6) # Reduce
font size of tick labels
    ax.set_xlabel(ax.get_xlabel(), fontsize=8) # Reduce font size of
```

```

x-axis labels
ax.set_ylabel(ax.get_ylabel(), fontsize=8) # Reduce font size of
y-axis labels
ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right',
fontsize=6) # Rotate x-axis labels and reduce font size

plt.tight_layout(pad=4.0) # Adjust spacing between subplots
plt.show()

```



```

perc = df.groupby(['customer_type', 'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
print(perc)

```

is_canceled	0	1	total perc	canc perc
customer_type				
Contract	2582	1262	0.035849	0.328304
Group	471	58	0.004933	0.109641
Transient	46833	33108	0.745517	0.414155
Transient-Party	16637	6278	0.213702	0.273969

Similarly, transient customer has the highest cancellation among customer types.

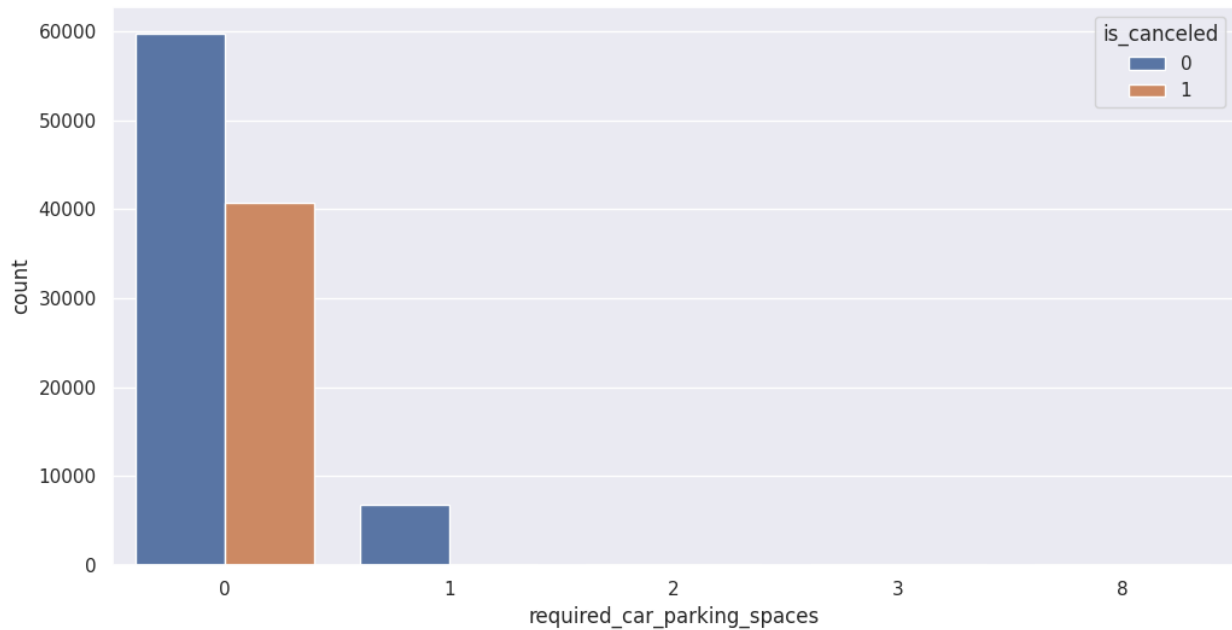
Parking space requirement

```

sns.countplot(x=df["required_car_parking_spaces"], hue=df["is_canceled"]
])

<Axes: xlabel='required_car_parking_spaces', ylabel='count'>

```



```
perc = df.groupby(['required_car_parking_spaces',
                  'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
print(perc)
```

is_canceled	0	1	total perc	canc perc
required_car_parking_spaces				
0	59732.0	40706.0	0.936668	0.405285
1	6759.0	NaN	NaN	NaN
2	27.0	NaN	NaN	NaN
3	3.0	NaN	NaN	NaN
8	2.0	NaN	NaN	NaN

93% of data belongs to people who did not require parking spaces. All the cancellations belong to this category only.

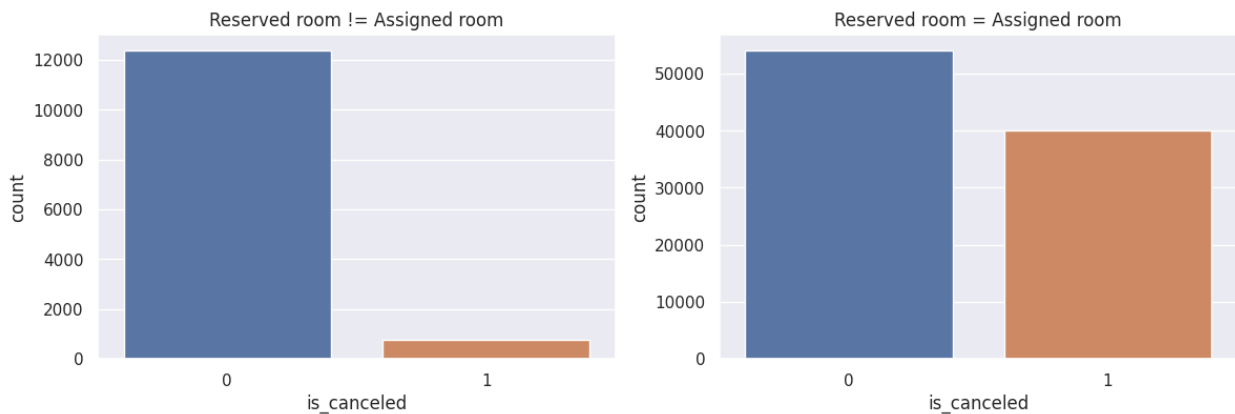
Room reservation and allotment

We can assume that people who were not assigned the same room that was reserved for them may cancel reservations. For this, we can take a separate dataset consisting only of data where reserved room type is not same as assigned room type and check.

```
rooms = df[(df["reserved_room_type"]) != (df["assigned_room_type"])]
rooms_no_change = df[(df["reserved_room_type"] ==
(df["assigned_room_type"])]

f, axes = plt.subplots(1,2,figsize=(14,4))
sns.countplot(x=rooms["is_canceled"],ax=axes[0]).set(title="Reserved
room != Assigned room")
```

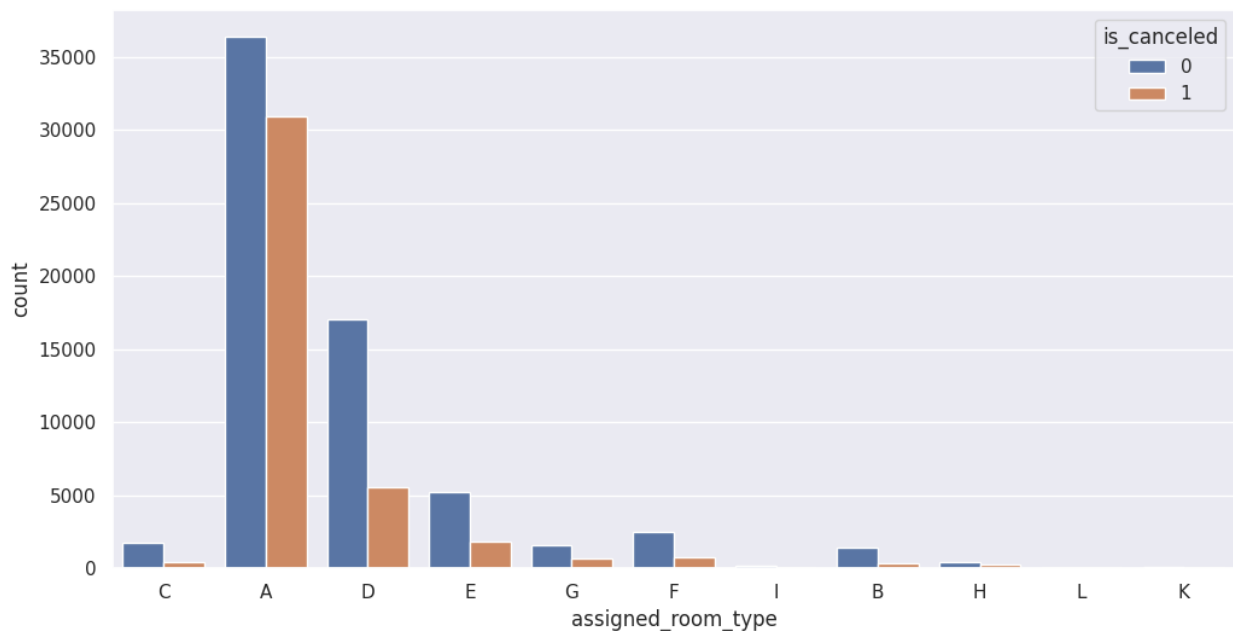
```
sns.countplot(x=rooms_no_change["is_canceled"],ax=axes[1]).set(title="Reserved room = Assigned room")
plt.show()
```



We can see that out of the data, the impact on cancellation is marginal even if assigned room is not same as reserved room. However, to reduce redundancy, we can combine it into a single column indicating if there is a change in reserved and assigned room.

Let us also check if assigned room type has any influence on the cancellation

```
sns.countplot(x=df["assigned_room_type"],hue=df["is_canceled"])
<Axes: xlabel='assigned_room_type', ylabel='count'>
```



We see that room A and D have majority guests and cancellations are also high.

Checking if change of room has any impact on cancellation

```
#Adding additional column to see if there is any change of room
df["room_change"] = 0

row_index = 0
for i in df["reserved_room_type"]:
    if i == df["assigned_room_type"][row_index]:
        df["room_change"][row_index] = 0
    else:
        df["room_change"][row_index] = 1
        row_index = row_index + 1

df.drop(columns = ["reserved_room_type", "assigned_room_type"],
inplace=True)
```

Previous cancellations and booking changes

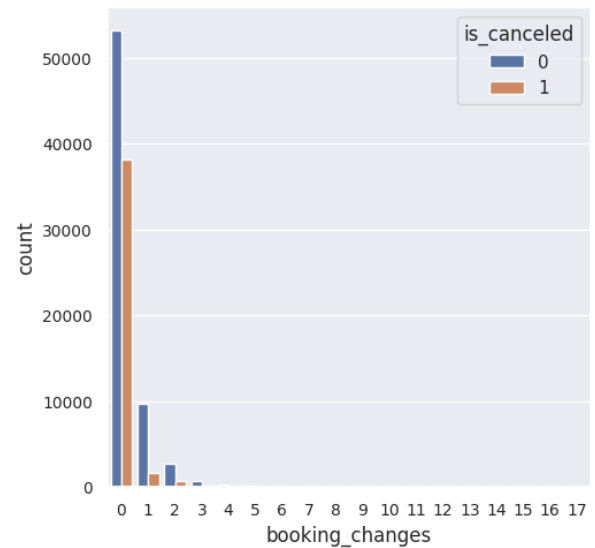
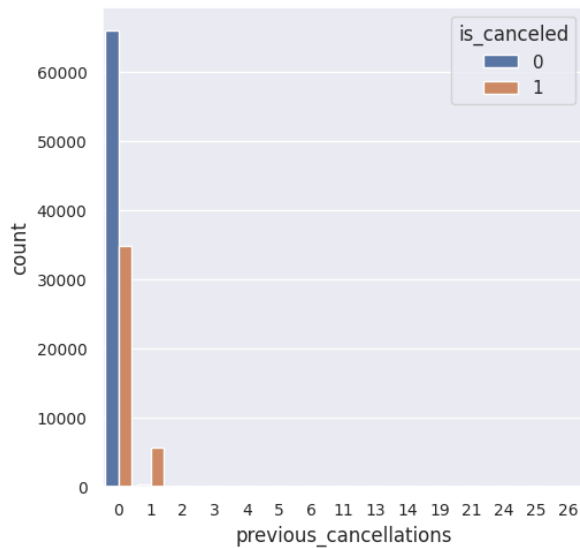
We can assume that people who have made previous cancellations have a higher likelihood to cancel this time. Similarly customers who have not canceled previous bookings will also not cancel this time.

```
sns.set(rc={"figure.figsize":(12,6)}) # Adjust figure size for better fit
f, axes = plt.subplots(1, 2, figsize=(12, 6)) # Adjust size for two plots side by side

# Plotting countplots
sns.countplot(x=df["previous_cancellations"], hue=df["is_canceled"],
ax=axes[0])
sns.countplot(x=df["booking_changes"], hue=df["is_canceled"],
ax=axes[1])

# Adjust font size of ticks and labels
for ax in axes:
    ax.tick_params(axis='both', which='major', labelsize=10) # Adjust font size of tick labels
    ax.set_xlabel(ax.get_xlabel(), fontsize=12) # Adjust font size of x-axis labels
    ax.set_ylabel(ax.get_ylabel(), fontsize=12) # Adjust font size of y-axis labels

plt.tight_layout(pad=4.0) # Adjust spacing between subplots
plt.show()
```

```
perc = df.groupby(['previous_cancellations',
'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
print(perc)
```

is_canceled	0	1	total perc	canc perc
previous_cancellations				
0	66005.0	34796.0	0.940054	0.345195
1	320.0	5683.0	0.055983	0.946693
2	74.0	37.0	0.001035	0.333333
3	45.0	20.0	0.000606	0.307692
4	24.0	7.0	0.000289	0.225806
5	17.0	2.0	0.000177	0.105263
6	15.0	7.0	0.000205	0.318182
11	22.0	10.0	0.000298	0.312500
13	1.0	11.0	0.000112	0.916667
14	NaN	14.0	NaN	NaN
19	NaN	19.0	NaN	NaN
21	NaN	1.0	NaN	NaN
24	NaN	48.0	NaN	NaN
25	NaN	25.0	NaN	NaN
26	NaN	26.0	NaN	NaN

```
perc = df.groupby(['booking_changes',
'is_canceled']).size().unstack(1)
perc["total perc"] = (perc[0] + perc[1])/len(df["is_canceled"])
perc["canc perc"] = (perc[1]/(perc[0] + perc[1]))
print(perc)
```

is_canceled	0	1	total perc	canc perc
booking_changes				
0	53159.0	38195.0	0.851952	0.418099

1	9649.0	1640.0	0.105279	0.145274
2	2636.0	650.0	0.030645	0.197809
3	660.0	123.0	0.007302	0.157088
4	251.0	57.0	0.002872	0.185065
5	79.0	14.0	0.000867	0.150538
6	34.0	16.0	0.000466	0.320000
7	21.0	3.0	0.000224	0.125000
8	10.0	4.0	0.000131	0.285714
9	7.0	1.0	0.000075	0.125000
10	5.0	1.0	0.000056	0.166667
11	1.0	NaN	NaN	NaN
12	1.0	NaN	NaN	NaN
13	5.0	NaN	NaN	NaN
14	1.0	1.0	0.000019	0.500000
15	2.0	NaN	NaN	NaN
16	NaN	1.0	NaN	NaN
17	2.0	NaN	NaN	NaN

- Here, we can see that cancellation percentage is high where previous cancellation is 1. Hence our first assumption is true.
- However, second assumption has failed. But that is because of the count of data. The column has maximum number of the same value.

Relationship with adr

adr - Average Daily Rate - It can be assumed that rooms with higher adr may be costlier. Hence, the chances of cancellation needs to be examined. As we do not have a feature for room price, we should check the relationship of adr with other factors also.

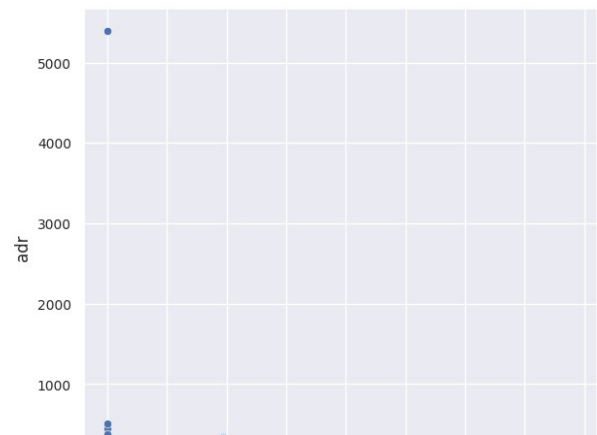
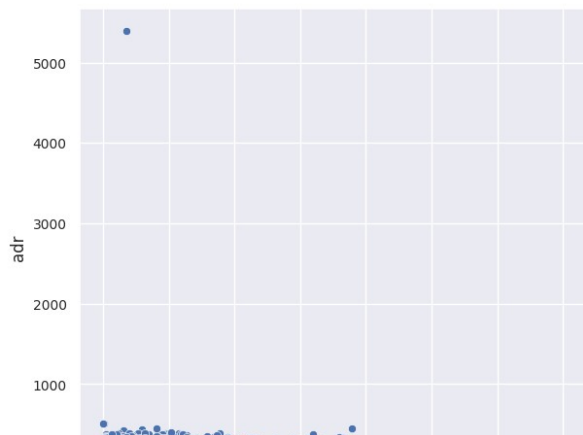
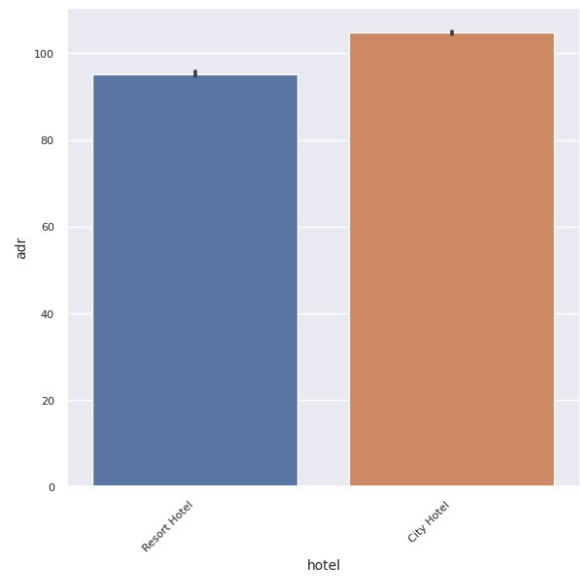
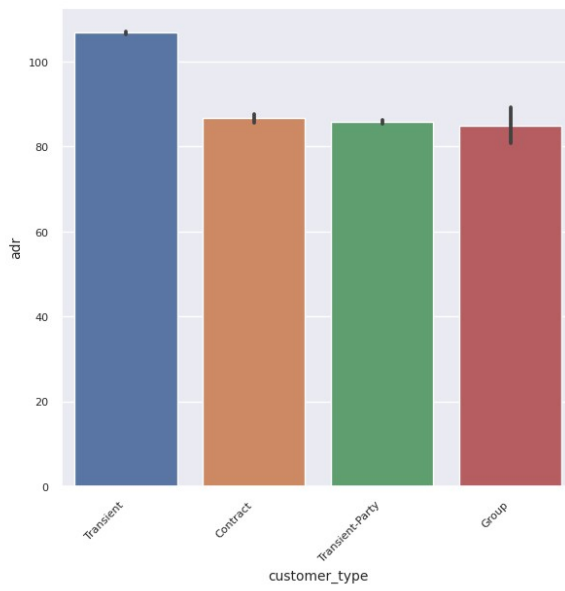
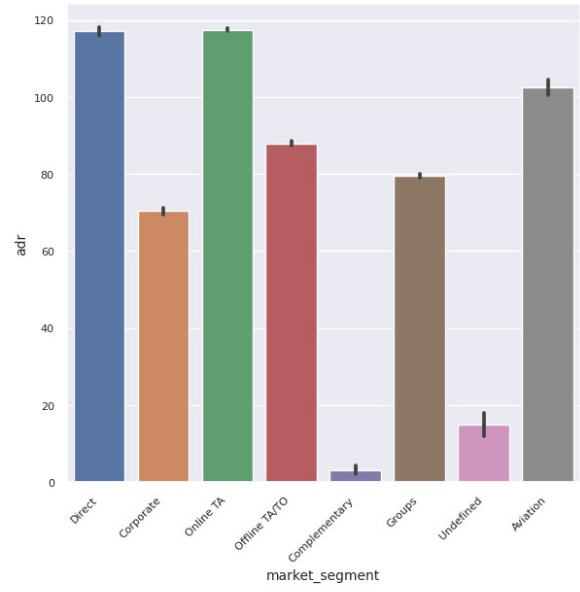
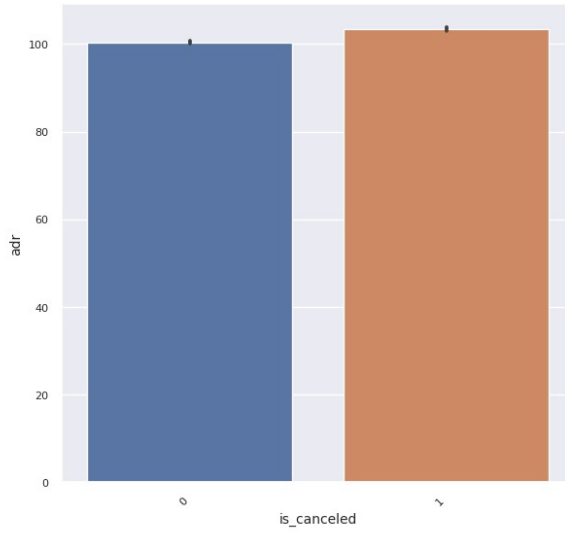
```
f, axes = plt.subplots(3, 2, figsize=(14, 20))

# Plotting barplots and scatterplots
sns.barplot(x=df["is_canceled"], y=df["adr"], ax=axes[0, 0])
sns.barplot(x=df["market_segment"], y=df["adr"], ax=axes[0, 1])
sns.barplot(x=df["customer_type"], y=df["adr"], ax=axes[1, 0])
sns.barplot(x=df["hotel"], y=df["adr"], ax=axes[1, 1])
sns.scatterplot(x=df["lead_time"], y=df["adr"], ax=axes[2, 0])
sns.scatterplot(x=df["days_in_waiting_list"], y=df["adr"], ax=axes[2, 1])

# Adjust font size and rotate labels for barplots
for ax in axes.flat:
    if ax in [axes[0, 0], axes[0, 1], axes[1, 0], axes[1, 1]]:
        ax.tick_params(axis='both', which='major', labelsize=8) #
Reduce font size of tick labels
        ax.set_xlabel(ax.get_xlabel(), fontsize=10) # Reduce font
size of x-axis labels
        ax.set_ylabel(ax.get_ylabel(), fontsize=10) # Reduce font
```

```
size of y-axis labels
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45,
ha='right', fontsize=8) # Rotate and reduce font size of x-axis
labels
    else:
        ax.tick_params(axis='both', which='major', labelsize=10) #
Adjust font size of tick labels for scatterplots
        ax.set_xlabel(ax.get_xlabel(), fontsize=12) # Adjust font
size of x-axis labels for scatterplots
        ax.set_ylabel(ax.get_ylabel(), fontsize=12) # Adjust font
size of y-axis labels for scatterplots

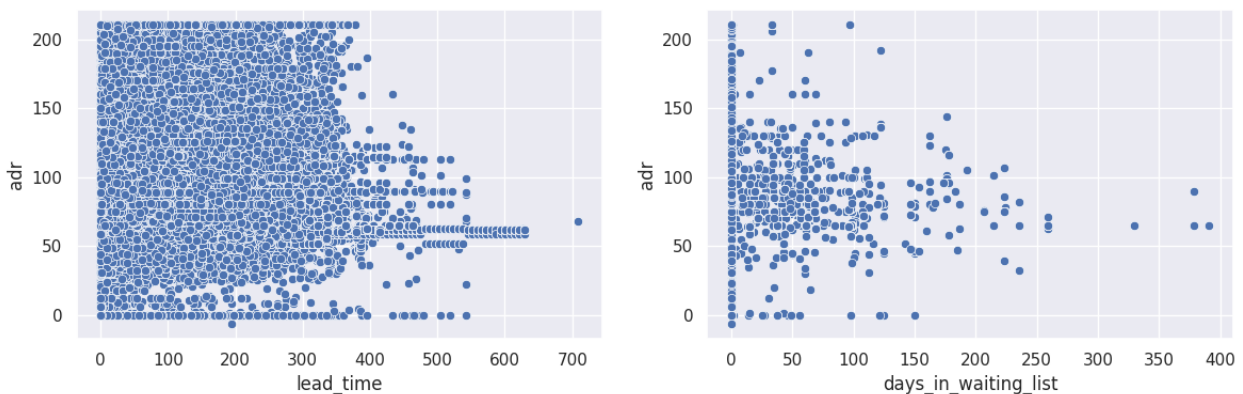
plt.tight_layout(pad=4.0) # Adjust spacing between subplots
plt.show()
```



- The mean adr for cancellation is slightly higher than non-cancellations.
- Direct and Online TA market segment have higher mean adr. The bookings in these categories are frequent and the likelihood of concessions is low.
- Transient customer type has higher mean adr. As transient customers are more frequent and seek shorter stays, they bring higher revenue.
- City hotels have higher adr.
- It appears that adr decreases as lead time and waiting list duration increases. But because of outliers, the pattern is unclear.

```
#Outlier treatment for waiting list and lead time columns
df1 = df.copy()
def outlier(data):
    Q1 = data.quantile(0.25)
    Q3 = data.quantile(0.75)
    IQR = Q3-Q1
    upper_bound = Q3 + 1.5*IQR
    lower_bound = Q1 - 1.5*IQR
    return data.clip(upper_bound, lower_bound)
for col in ['adr']:
    df1[col] = outlier(df1[col])

f, axes = plt.subplots(1, 2, figsize=(14, 4))
sns.scatterplot(x=df1["lead_time"], y=df1["adr"], ax=axes[0])
sns.scatterplot(x=df1["days_in_waiting_list"], y=df1["adr"], ax=axes[1])
plt.show()
```



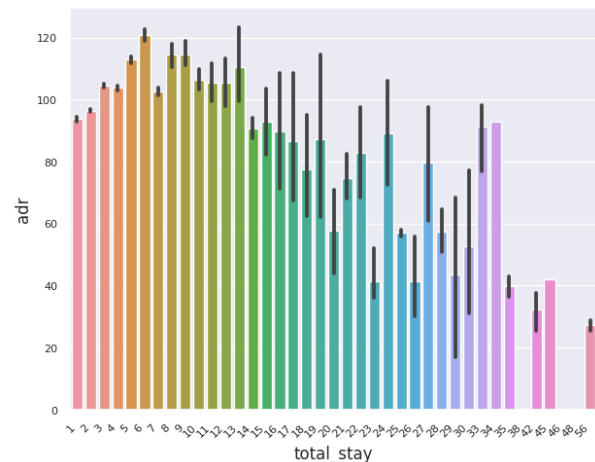
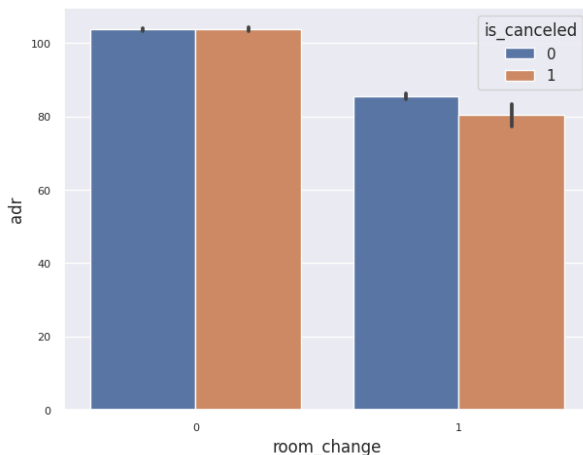
There seems to be no relationship between lead time and adr or waiting list and adr after treating outliers.

```
# Set figure size and create subplots
f, axes = plt.subplots(1, 2, figsize=(14, 6))

# Plotting the first barplot
sns.barplot(x=df["room_change"], y=df["adr"], hue=df["is_canceled"],
ax=axes[0])
axes[0].tick_params(axis='both', which='major', labelsize=8) #
Further reduce font size of tick labels
```

```
# Plotting the second barplot
sns.barplot(x=df["total_stay"], y=df["adr"], ax=axes[1])
axes[1].tick_params(axis='both', which='major', labelsize=8) #
Further reduce font size of tick labels
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=45,
ha='right', fontsize=8) # Rotate and reduce font size of x-axis
labels

plt.tight_layout(pad=4.0) # Adjust spacing between subplots
plt.show()
```



- Room changes have lower mean adr than no room changes. And cancellation after room changes have even lesser adr. This can indicate that hotels incur a loss due to room changes and further cancellation.
- As the total stay duration increases, the mean adr increases and then drops with further increase in duration of stay.

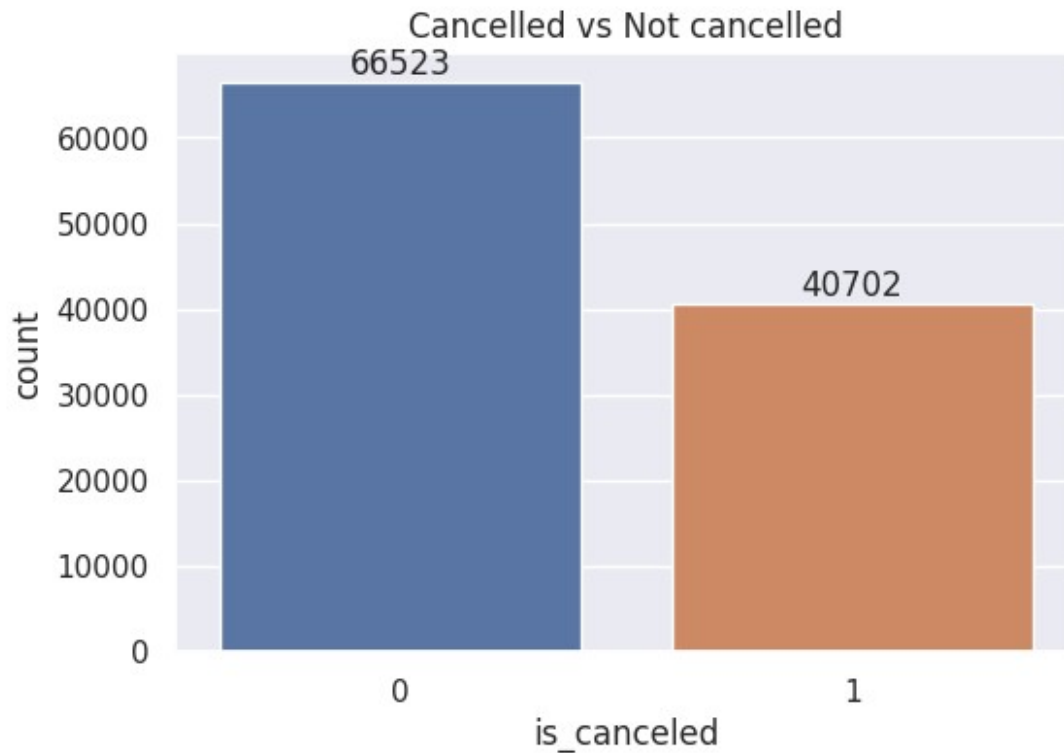
```
#Dropping the remaining records with missing values
df = df.dropna(subset = ["kids"])

df = df.reset_index()
df = df.drop(columns=["index"])
```

Checking the imbalance in data after preliminary preprocessing

```
plt.figure(figsize=(6,4))

ax = sns.countplot(x=df["is_canceled"])
for i in ax.containers:
    ax.bar_label(i)
ax.set(title="Cancelled vs Not cancelled")
plt.show()
```



The distribution remains similar after all the eliminations. We will not consider the data as imbalanced because there is a significant amount of positive class.

Correlations

First we encode categorical columns to check correlation and we do this by creating a separate dataset so as not to disturb our original dataset

```
corr_df = df.copy()

corr_df["month"] =
corr_df["arrival_date_month"].map({"January":1,"February":2,"March":3,
"April":4,"May":5,"June":6,"July":7,"August":8,
"September":9,"October":10,"November":11,"December":12})
corr_df["year"] = corr_df["arrival_date_year"].astype(str)

corr_df["year"] = corr_df["year"].map({"2015":1,"2016":2,"2017":3})

corr_df = corr_df.drop(columns =
["arrival_date_month","arrival_date_year"])

cat_cols = corr_df.select_dtypes(include='object').columns

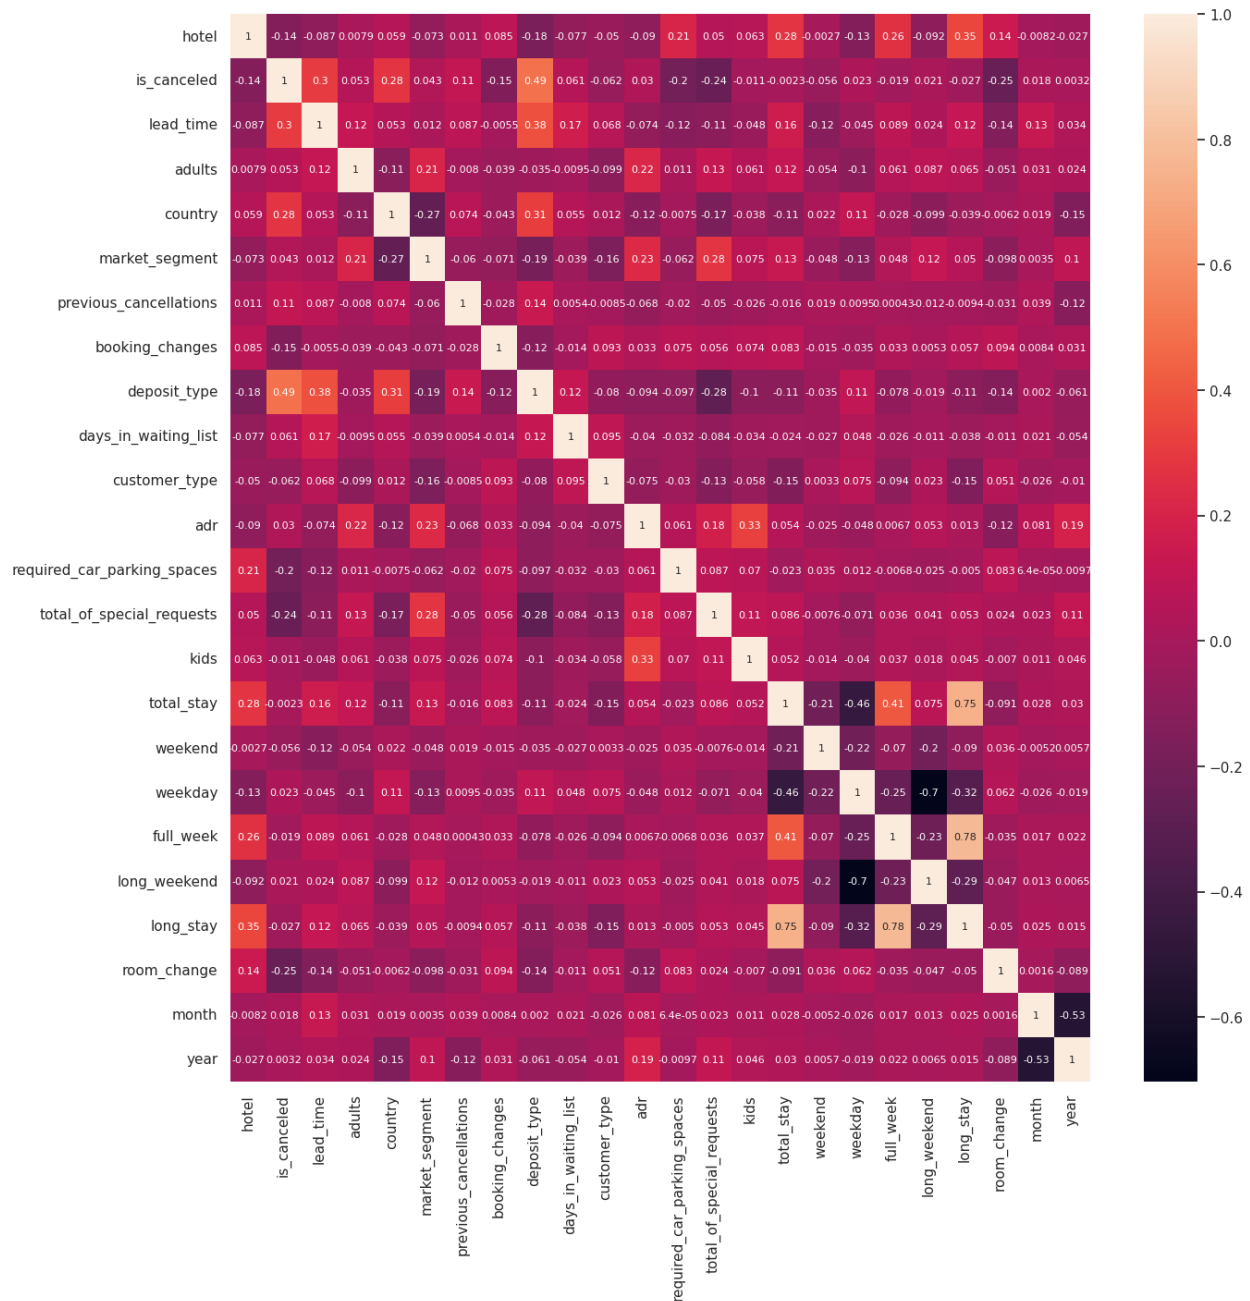
#Encoding all categorical columns
from sklearn.preprocessing import LabelEncoder
obj = LabelEncoder()
```

```

for i in cat_cols:
    corr_df[i] = obj.fit_transform(corr_df[i])

plt.figure(figsize=(15, 15))
sns.heatmap(corr_df.corr(), annot=True, annot_kws={"size": 8})
plt.show()

```



Two columns have a moderate to high correlation which are total stay, full week and long stay. We will drop the redundant column.

Dropping correlated columns from original dataset


```
df = df.drop(columns = ["long_stay"])
```

Converting categorical columns

Now we need to convert categorical columns into numerical

There are ordinal columns: Month and year

The rest can be kept as nominal.

```
from sklearn.preprocessing import LabelEncoder

obj = LabelEncoder()
df["country"] = obj.fit_transform(df["country"])
#df["assigned_room_type"] =
obj.fit_transform(df["assigned_room_type"])

#Manual encoding of ordinal column
df["month"] =
df["arrival_date_month"].map({"January":1,"February":2,"March":3,"April":4,"May":5,"June":6,"July":7,"August":8,
"September":9,"October":10,"November":11,"December":12})

df["year"] = df["arrival_date_year"].astype(str)

df["year"] = df["year"].map({"2015":1,"2016":2,"2017":3})

df = df.drop(columns = ["arrival_date_month","arrival_date_year"])

dummy_cols = ["hotel","market_segment","deposit_type","customer_type"]

#One-hot encoding
df = pd.get_dummies(df, columns = dummy_cols)
```

Standardization and Feature importances

```
X = df.drop(columns = ["is_canceled"])
y = df["is_canceled"]

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,
random_state=23, stratify=y)

#Calculating feature importances
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(X_train,y_train)

RandomForestClassifier()
```

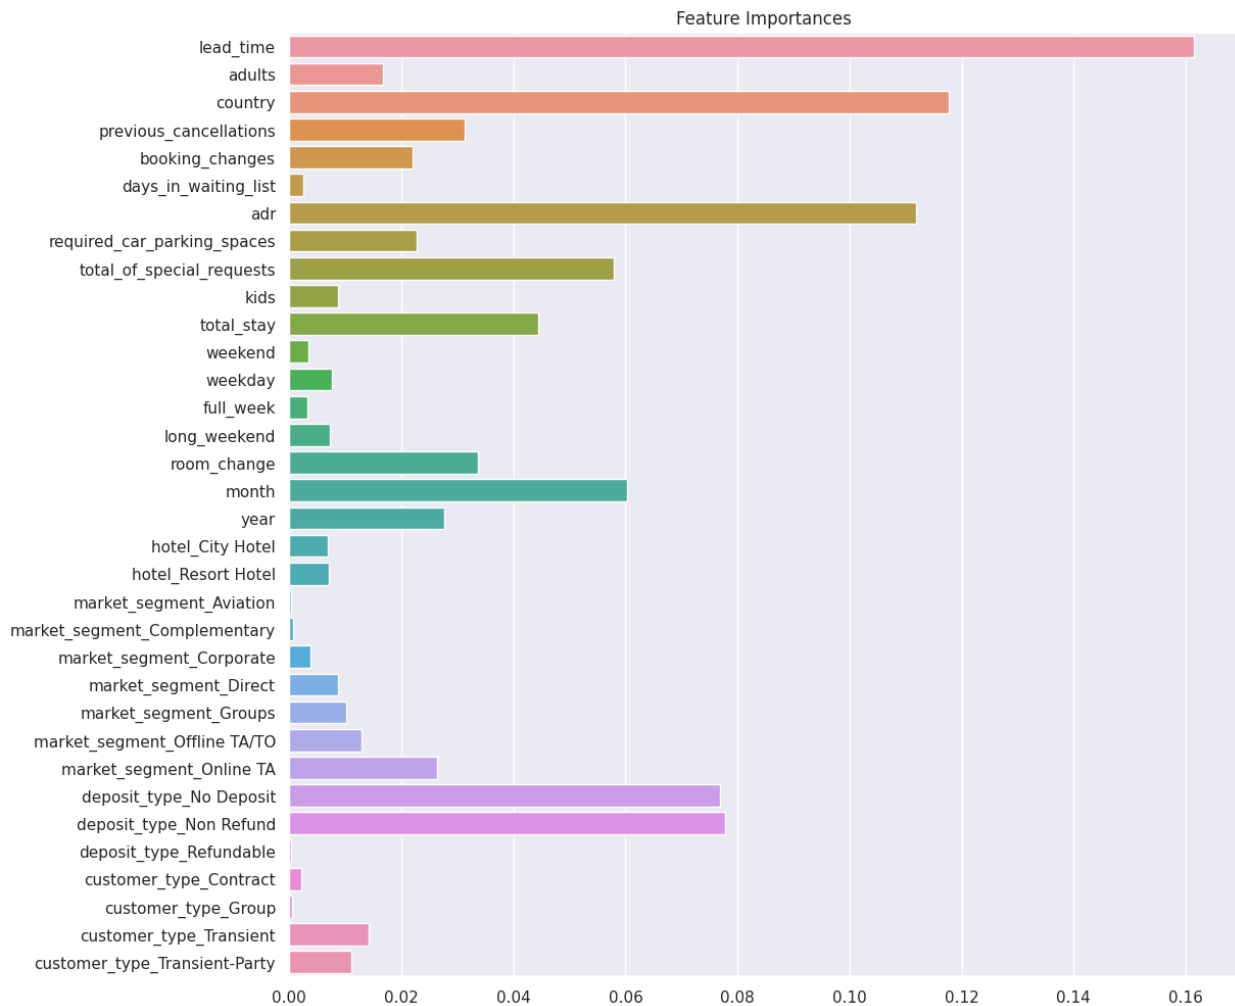
```

fi=rfc.feature_importances_
xc = X_train.columns

sns.set(rc={"figure.figsize":(12,12)})
sns.barplot(y=xc,x=fi).set(title="Feature Importances")

[Text(0.5, 1.0, 'Feature Importances')]

```



```

#Dropping less important features after finding Feature Importances
X = X.drop(columns =
["market_segment_Aviation","market_segment_Complementary","weekend","full_week",
"days_in_waiting_list",

"deposit_type_Refundable","customer_type_Group","customer_type_Contract"

])

#Re-splitting data into train and test sets after dropping less important features

```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,
random_state=41, stratify=y)
```

Normalizing Training Data

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
sc = scaler.fit_transform(X_train)

cols = X_train.columns
X_train = pd.DataFrame(sc, columns = cols)
```

Training different models one by one and testing their metrics

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()

from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(solver='liblinear')

from sklearn.naive_bayes import BernoulliNB
nb = BernoulliNB()

from sklearn.ensemble import GradientBoostingClassifier
gbc = GradientBoostingClassifier()

model_rfc = rfc.fit(X_train,y_train)
model_lr = lr.fit(X_train,y_train)
model_nb = nb.fit(X_train,y_train)
model_gbc = gbc.fit(X_train,y_train)
```

Normalizing Test Data

```
sc = scaler.transform(X_test)
cols = X_test.columns
X_test = pd.DataFrame(sc, columns = cols)
```

RandomForestClassifier

```
y_pred_test_rfc = model_rfc.predict(X_test)
y_pred_train_rfc = model_rfc.predict(X_train)

from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
print(f"Metrics for test set")
result_cm = confusion_matrix(y_test, y_pred_test_rfc)
print(f"Confusion matrix {model_rfc}: \n", result_cm)
result_cr = classification_report(y_test, y_pred_test_rfc)
print(f"Classification Report {model_rfc}: \n", result_cr)
```

```

acc_score_rfc = accuracy_score(y_test,y_pred_test_rfc)
print(f"Accuracy Score: for {model_rfc} \n", acc_score_rfc)

print(f"Metrics for training set")
result_cm = confusion_matrix(y_train, y_pred_train_rfc)
print(f"Confusion matrix {model_rfc}: \n", result_cm)
result_cr = classification_report(y_train, y_pred_train_rfc)
print(f"Classification Report {model_rfc}: \n", result_cr)
acc_score = accuracy_score(y_train,y_pred_train_rfc)
print(f"Accuracy Score: for {model_rfc} \n", acc_score)

```

Metrics for test set

Confusion matrix RandomForestClassifier():

```

[[18526  1431]
 [ 2163 10048]]

```

Classification Report RandomForestClassifier():

	precision	recall	f1-score	support
0	0.90	0.93	0.91	19957
1	0.88	0.82	0.85	12211
accuracy			0.89	32168
macro avg	0.89	0.88	0.88	32168
weighted avg	0.89	0.89	0.89	32168

Accuracy Score: for RandomForestClassifier()

0.8882740611788112

Metrics for training set

Confusion matrix RandomForestClassifier():

```

[[46434   132]
 [  211 28280]]

```

Classification Report RandomForestClassifier():

	precision	recall	f1-score	support
0	1.00	1.00	1.00	46566
1	1.00	0.99	0.99	28491
accuracy			1.00	75057
macro avg	1.00	0.99	1.00	75057
weighted avg	1.00	1.00	1.00	75057

Accuracy Score: for RandomForestClassifier()

0.9954301397604487

```

from sklearn.metrics import roc_auc_score

```

```

auc_rfc = np.round(roc_auc_score(y_test,y_pred_test_rfc),3)

```

```

auc_rfc

```

0.876

- For Random Forest Classifier we can see that the test set has an accuracy score of around 88% while the training set has an accuracy score of around 99%. The difference in accuracy score between training and testing set is around 11% which indicates overfitting.
- The ROC-AUC score is around 87% which is close to the accuracy score of the test set. This means that the model has the ability to classify 87% of the data correctly.

Logistic Regression

```
y_pred_test_lr = model_lr.predict(X_test)
y_pred_train_lr = model_lr.predict(X_train)

from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
print(f"Metrics for test set")
result_cm = confusion_matrix(y_test, y_pred_test_lr)
print(f"Confusion matrix {model_lr}: \n", result_cm)
result_cr = classification_report(y_test, y_pred_test_lr)
print(f"Classification Report {model_lr}: \n", result_cr)
acc_score_lr = accuracy_score(y_test, y_pred_test_lr)
print(f"Accuracy Score: for {model_lr} \n", acc_score_lr)
```

```
print(f"Metrics for training set")
result_cm = confusion_matrix(y_train, y_pred_train_lr)
print(f"Confusion matrix {model_lr}: \n", result_cm)
result_cr = classification_report(y_train, y_pred_train_lr)
print(f"Classification Report {model_lr}: \n", result_cr)
acc_score = accuracy_score(y_train, y_pred_train_lr)
print(f"Accuracy Score: for {model_lr} \n", acc_score)
```

Metrics for test set

Confusion matrix LogisticRegression(solver='liblinear'):

```
[[18148 1809]
 [ 4627 7584]]
```

Classification Report LogisticRegression(solver='liblinear'):

	precision	recall	f1-score	support
0	0.80	0.91	0.85	19957
1	0.81	0.62	0.70	12211
accuracy			0.80	32168
macro avg	0.80	0.77	0.78	32168
weighted avg	0.80	0.80	0.79	32168

Accuracy Score: for LogisticRegression(solver='liblinear')
0.7999253916936085

Metrics for training set

Confusion matrix LogisticRegression(solver='liblinear'):

```
[[42394 4172]
 [10679 17812]]
```

Classification Report LogisticRegression(solver='liblinear'):

	precision	recall	f1-score	support
0	0.80	0.91	0.85	46566
1	0.81	0.63	0.71	28491
accuracy			0.80	75057
macro avg	0.80	0.77	0.78	75057
weighted avg	0.80	0.80	0.80	75057

Accuracy Score: for LogisticRegression(solver='liblinear')
0.8021370425143558

```
from sklearn.metrics import roc_auc_score
auc_lr = np.round(roc_auc_score(y_test,y_pred_test_lr),3)
auc_lr
```

0.765

For Logistic Regression

- Test Accuracy: 80%
- Train Accuracy : 80%
- AUC score : 0.765

BernoulliNB

```
y_pred_test_bnb = model_nb.predict(X_test)
y_pred_train_bnb = model_nb.predict(X_train)

from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
print(f"Metrics for test set")
result_cm = confusion_matrix(y_test, y_pred_test_bnb)
print(f"Confusion matrix {model_nb}: \n", result_cm)
result_cr = classification_report(y_test, y_pred_test_bnb)
print(f"Classification Report {model_nb}: \n", result_cr)
acc_score_bnb = accuracy_score(y_test,y_pred_test_bnb)
print(f"Accuracy Score: for {model_nb} \n", acc_score_bnb)

print(f"Metrics for training set")
result_cm = confusion_matrix(y_train, y_pred_train_bnb)
print(f"Confusion matrix {model_nb}: \n", result_cm)
result_cr = classification_report(y_train, y_pred_train_bnb)
print(f"Classification Report {model_nb}: \n", result_cr)
acc_score = accuracy_score(y_train,y_pred_train_bnb)
print(f"Accuracy Score: for {model_nb} \n", acc_score)
```

Metrics for test set
Confusion matrix BernoulliNB():
[[18400 1557]
[5598 6613]]

```

Classification Report BernoulliNB():
              precision    recall  f1-score   support

     0       0.77       0.92       0.84     19957
     1       0.81       0.54       0.65     12211

 accuracy          0.78     32168
 macro avg       0.79     32168
 weighted avg    0.78     32168

```

```

Accuracy Score: for BernoulliNB()
0.7775739865705048

```

Metrics for training set

```

Confusion matrix BernoulliNB():
[[42862  3704]
 [13009 15482]]

```

```

Classification Report BernoulliNB():
              precision    recall  f1-score   support

     0       0.77       0.92       0.84     46566
     1       0.81       0.54       0.65     28491

 accuracy          0.78     75057
 macro avg       0.79     75057
 weighted avg    0.78     75057

```

```

Accuracy Score: for BernoulliNB()
0.7773292297853631

```

```

from sklearn.metrics import roc_auc_score
auc_bnb = np.roun(roc_auc_score(y_test,y_pred_test_bnb),3)
auc_bnb
0.732

```

For Bernoulli Naive Bayes model,

- Test Accuracy: 77%
- Train Accuracy : 77%
- AUC score : 0.732

Gradient Boosting Classifier

```

y_pred_test_gbc = model_gbc.predict(X_test)
y_pred_train_gbc = model_gbc.predict(X_train)

from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
print(f"Metrics for test set")
result_cm = confusion_matrix(y_test, y_pred_test_gbc)
print(f"Confusion matrix {model_gbc}: \n", result_cm)

```

```

result_cr = classification_report(y_test, y_pred_test_gbc)
print(f"Classification Report {model_gbc}: \n", result_cr)
acc_score_gbc = accuracy_score(y_test, y_pred_test_gbc)
print(f"Accuracy Score: for {model_gbc} \n", acc_score_gbc)

print(f"Metrics for training set")
result_cm = confusion_matrix(y_train, y_pred_train_gbc)
print(f"Confusion matrix {model_gbc}: \n", result_cm)
result_cr = classification_report(y_train, y_pred_train_gbc)
print(f"Classification Report {model_gbc}: \n", result_cr)
acc_score = accuracy_score(y_train, y_pred_train_gbc)
print(f"Accuracy Score: for {model_gbc} \n", acc_score)

```

Metrics for test set

Confusion matrix GradientBoostingClassifier():

```

[[18041  1916]
 [ 2888  9323]]

```

Classification Report GradientBoostingClassifier():

	precision	recall	f1-score	support
0	0.86	0.90	0.88	19957
1	0.83	0.76	0.80	12211
accuracy			0.85	32168
macro avg	0.85	0.83	0.84	32168
weighted avg	0.85	0.85	0.85	32168

Accuracy Score: for GradientBoostingClassifier()

0.8506590400397911

Metrics for training set

Confusion matrix GradientBoostingClassifier():

```

[[42174  4392]
 [ 6945 21546]]

```

Classification Report GradientBoostingClassifier():

	precision	recall	f1-score	support
0	0.86	0.91	0.88	46566
1	0.83	0.76	0.79	28491
accuracy			0.85	75057
macro avg	0.84	0.83	0.84	75057
weighted avg	0.85	0.85	0.85	75057

Accuracy Score: for GradientBoostingClassifier()

0.8489547943562892

```

from sklearn.metrics import roc_auc_score

```

```

auc_gbc = np.round(roc_auc_score(y_test, y_pred_test_gbc), 3)

```

```

auc_gbc

```

0.834

For Gradient Boosting Classifier,

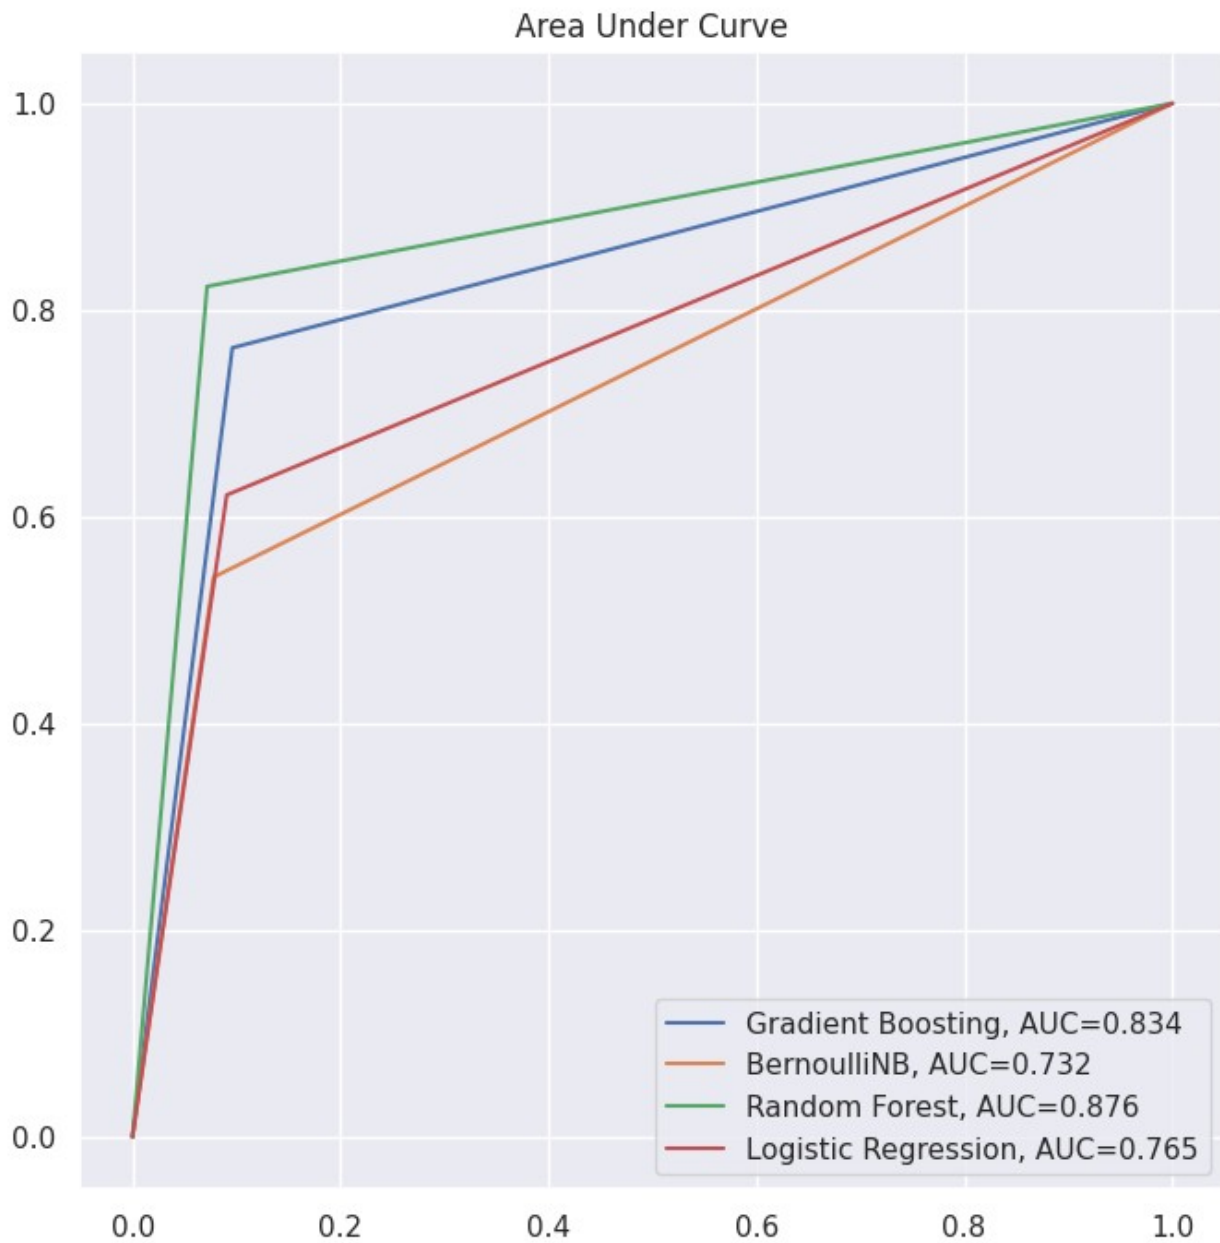
- Test Accuracy: 85%
- Train Accuracy : 85%
- AUC score : 0.834

Visualizing performance of all models

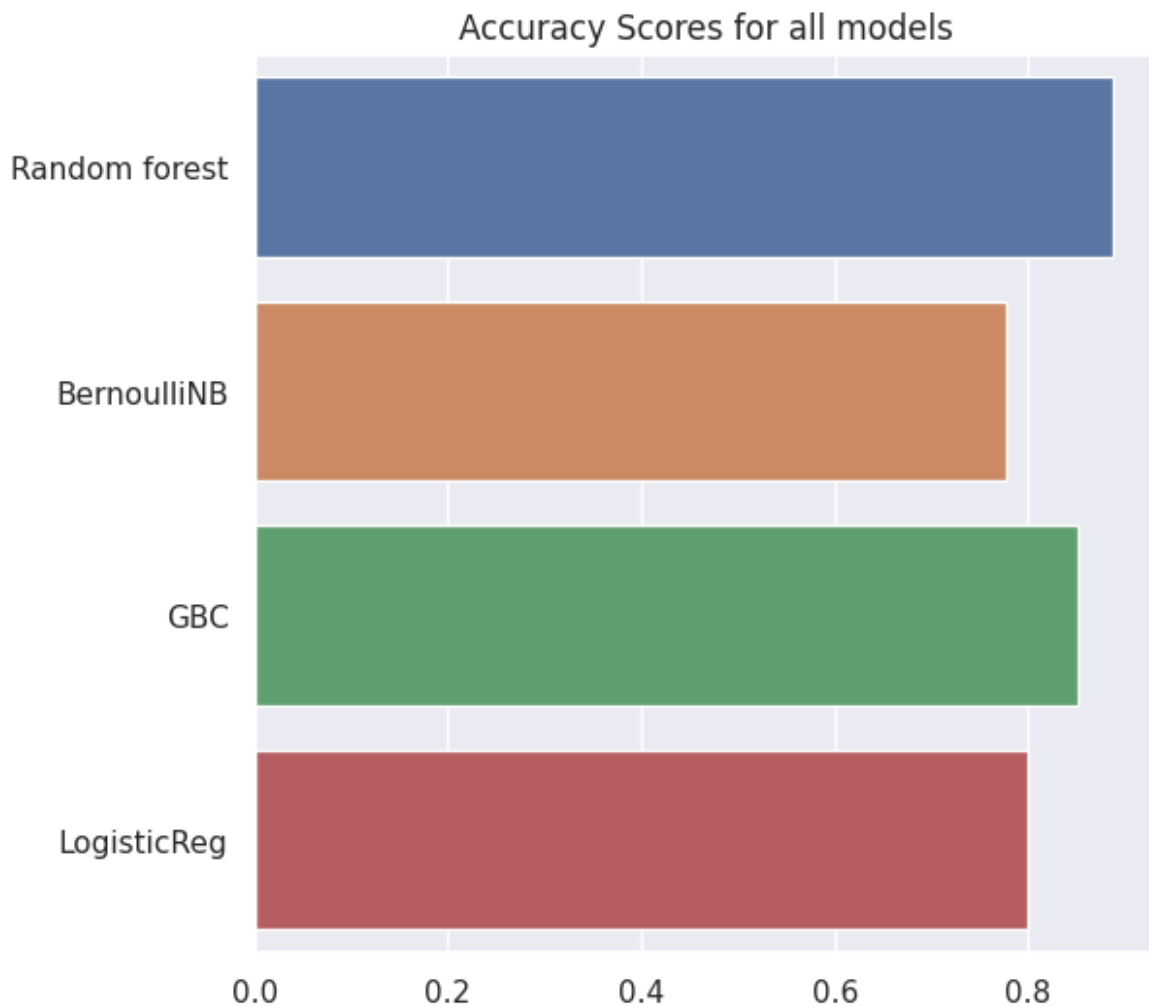
```
acc = [acc_score_rfc, acc_score_bnb, acc_score_gbc, acc_score_lr]
m = ["Random forest", "BernoulliNB", "GBC", "LogisticReg"]

from sklearn import metrics
plt.figure(figsize=(8,8))
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_test_gbc)
plt.plot(fpr, tpr, label="Gradient Boosting, AUC="+str(auc_gbc))
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_test_bnb)
plt.plot(fpr, tpr, label="BernoulliNB, AUC="+str(auc_bnb))
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_test_rfc)
plt.plot(fpr, tpr, label="Random Forest, AUC="+str(auc_rfc))
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_test_lr)
plt.plot(fpr, tpr, label="Logistic Regression, AUC="+str(auc_lr))
plt.legend()
plt.title("Area Under Curve")
```

```
Text(0.5, 1.0, 'Area Under Curve')
```



```
plt.figure(figsize=(6,6))
sns.barplot(x=acc,y=m)
plt.title("Accuracy Scores for all models")
plt.show()
```



Random Forest Classifier model has highest testing and training accuracies and also AUC score of 0.876. However, there is 11% difference between training and testing accuracy scores.

We choose Random Forest Classifier as it has highest accuracy and AUC score. We will perform K-Fold Cross-validation on training set to get the average training accuracy.

K-Fold Cross-validation of model

```
from sklearn.model_selection import cross_val_score

train_scores =
cross_val_score(model_rfc,X_train,y_train,cv=10,scoring='accuracy')
test_scores =
cross_val_score(model_rfc,X_test,y_test,cv=10,scoring='accuracy')

train_avg = train_scores.mean()
test_avg = test_scores.mean()

print(f"Training Average : {train_avg}")
print(f"Training scores : {train_scores}")
```

```

print("-"*100)
print(f"Test Average : {test_avg}")
print(f"Test scores : {test_scores}")

Training Average : 0.8856202819223636
Training scores : [0.88462563 0.88915534 0.88169464 0.89155342
0.88569145 0.88489209
0.88022915 0.88367755 0.8898068 0.88487675]
-----
-----
Test Average : 0.8747510504672782
Test scores : [0.87410631 0.87410631 0.88343177 0.86167237 0.87410631
0.87876904
0.87783649 0.88218837 0.86629353 0.875 ]

```

Average performance of Random Forest Classifier model is quite good and it has average accuracy of 87% for test set and 88% for training set.

Conclusion

In summary, the Random Forest Classifier yielded the best results for our training, testing, and cross-validation datasets. With an accuracy of 87%, it is likely that the model will classify new, unseen data with an accuracy between 75% and 85%.